

CS 33

Multithreaded Programming V

The source code used in this lecture is available on the course web page.

Implementing Mutexes (1)

- **Strategy**
 - make the usual case (no waiting) very fast
 - can afford to take more time for the other case (waiting for the mutex)

Implementing Mutexes (2)

- **Mutex has three states**
 - unlocked
 - locked, no waiters
 - locked, waiting threads
- **Locking the mutex**
 - use `cmpxchg` (with lock prefix)
 - if unlocked, lock it and we're done
 - » state changed to locked, no waiters
 - otherwise, make “futex” system call to wait till it's unlocked
 - » state changed to locked, waiting threads

The actual solution involves some complications not apparent here — see <http://people.redhat.com/drepper/futex.pdf> for details.

Implementing Mutexes (3)

- **Unlocking the mutex**
 - if locked, but no waiters
 - » state changed to unlocked
 - if locked, but waiting threads
 - » futex system call made to wake up a waiting thread

Memory Allocation

- Multiple threads
 - One heap
- Bottleneck?**

In a naïve multithreaded implementation of malloc/free, there is one mutex protecting the heap, resulting in a bottleneck.

Solution 1

- **Divvy up the heap among the threads**
 - each thread has its own heap
 - no mutexes required
 - no bottleneck
- **How much heap does each thread get?**

Solution 2

- **Global heap plus per-thread heaps**
 - threads pull storage from global heap
 - freed storage goes to per-thread heap
 - » unless things are imbalanced
 - then thread moves storage back to global heap
 - mutex on only the global heap
- **What if one thread allocates and another frees storage?**

Solution 3

- **Multiple “arenas”**
 - each with its own mutex
 - thread allocates from the first one it can find whose mutex was unlocked
 - » if none, then creates new one
 - deallocations go back to original arena

Malloc/Free Implementations

- **ptmalloc**
 - based on solution 3
 - in glibc (i.e., used by default)
- **tcmalloc**
 - based on solution 2
 - from Google
- **Which is best?**

Test Program

```
const unsigned int N=64, nthreads=32, iters=10000000;
int main() {
    void *tfunc(void *);
    pthread_t thread[nthreads];
    for (int i=0; i<nthreads; i++) {
        pthread_create(&thread[i], 0, tfunc, (void *)i);
        pthread_detach(thread[i]);
    }
    pthread_exit(0);
}
void *tfunc(void *arg) {
    long i;
    for (i=0; i<iters; i++) {
        long *p = (long *)malloc(sizeof(long)*((i%N)+1));
        free(p);
    }
    return 0;
}
```

Compiling It ...

```
% gcc -o ptalloc alloc.cc -lpthread  
% gcc -o tcalloc alloc.cc -lpthread -ltcmalloc
```

Running It ...

```
$ time ./ptalloc
real    0m5.142s
user    0m20.501s
sys     0m0.024s
$ time ./tcalloc
real    0m1.889s
user    0m7.492s
sys     0m0.008s
```

The code was run on an Intel(R) Core(TM)2 Quad CPU Q6600 @ 2.40GHz.

What's Going On?

```
$ strace -c -f ./ptalloc
```

```
...
% time      seconds  usecs/call   calls   errors syscall
-----
100.00      0.040002         13     3007      520 futex
...
```

```
$ strace -c -f ./tcalloc
```

```
...
% time      seconds  usecs/call   calls   errors syscall
-----
...
 0.00      0.000000         0         59      13 futex
...
```

Test Program 2, part 1

```
#define N 64
#define npairs 16
#define allocsPerIter 1024
const long iters = 8*1024*1024/allocsPerIter;
#define BufSize 10240
typedef struct buffer {
    int *buf[BufSize];
    unsigned int nextin;
    unsigned int nextout;
    sem_t empty;
    sem_t occupied;
    pthread_t pthread;
    pthread_t cthread;
} buffer_t;
```

This program creates pairs of threads: one thread allocates storage, the other deallocates storage. They communicate using producer-consumer communication.

Test Program 2, part 2

```
int main() {
    long i;
    buffer_t b[npairs];
    for (i=0; i<npairs; i++) {
        b[i].nextin = 0;
        b[i].nextout = 0;
        sem_init(&b[i].empty, 0, BufSize/allocsPerIter);
        sem_init(&b[i].occupied, 0, 0);
        pthread_create(&b[i].pthread, 0, prod, &b[i]);
        pthread_create(&b[i].cthread, 0, cons, &b[i]);
    }
    for (i=0; i<npairs; i++) {
        pthread_join(b[i].pthread, 0);
        pthread_join(b[i].cthread, 0);
    }
    return 0;
}
```

The main routine creates *npairs* (16) of communicating pairs of threads.

Test Program 2, part 3

```
void *prod(void *arg) {
    long i, j;
    buffer_t *b = (buffer_t *)arg;
    for (i = 0; i<iters; i++) {
        sem_wait(&b->empty);
        for (j = 0; j<allocsPerIter; j++) {
            b->buf[b->nextin] = malloc(sizeof(int)*((j%N)+1));
            if (++b->nextin >= BufSize)
                b->nextin = 0;
        }
        sem_post(&b->occupied);
    }
    return 0;
}
```

To reduce the number of calls to *sem_wait* and *sem_post*, at each iteration the thread calls *new allocsPerIter* (1024) times.

Test Program 2, part 4

```
void *cons(void *arg) {
    long i, j;
    buffer_t *b = (buffer_t *)arg;
    for (i = 0; i<iters; i++) {
        sem_wait(&b->occupied);
        for (j = 0; j<allocsPerIter; j++) {
            free(b->buf[b->nextout]);
            if (++b->nextout >= BufSize)
                b->nextout = 0;
        }
        sem_post(&b->empty);
    }
    return 0;
}
```

Running It ...

```
$ time ./ptalloc2
real    0m1.087s
user    0m3.744s
sys     0m0.204s
$ time ./talloc2
real    0m3.535s
user    0m11.361s
sys     0m2.112s
```

The code was run on a SunLab machine (an Intel(R) Core(TM)2 Quad CPU Q6600 @ 2.40GHz).

What's Going On?

```
$ strace -c -f ./ptalloc2
```

```
...
```

% time	seconds	usecs/call	calls	errors	syscall
--------	---------	------------	-------	--------	---------

94.96	2.347314	44	53653	14030	futex
-------	----------	----	-------	-------	-------

```
...
```

```
$ strace -c -f ./tccalloc2
```

```
...
```

% time	seconds	usecs/call	calls	errors	syscall
--------	---------	------------	-------	--------	---------

93.86	6.604632	36	185731	45222	futex
-------	----------	----	--------	-------	-------

```
...
```