

Welcome to CS110: Principles of Computer Systems

■ I'm Jerry Cain (jerry@cs.stanford.edu)

- Chemistry undergrad MIT, originally Ph.D. in Chemistry here, defected to CS in 1993
- Lecturer in CS, teaching CS106B, CS106X, CS107, CS110, CS143
- Taught CS110 for the first time in Spring, 2013
 - Leveraged much of Mendel Roseblum's CS110 materials from prior offerings
 - Introduced some of my own material since then, and will be introducing even more this time
 - CS110 is an evolving system, but I don't expect you to notice one bit
- Started working at [Facebook](#) in 2008, have worked on Open Graph team entire time (currently at one day a week)
- Learned web programming, PHP, CSS, JavaScript. Old CS107 student (class of 2004) is my manager [#karma](#)
- Have grown to understand and appreciate large systems much better as a result of working there

■ Staff and Students

- 205 enrolled students as of January 4th
- Each of you should know C and C++ reasonably well so that you can...
 - write moderately complex programs
 - read and understand portions of medium and even large code bases
 - calmly trace memory diagrams and always win
- Each of you should be fluent with Unix, **make**, **valgrind**, and **make** to the extent they're covered in CS107.
- 7 CA's at the moment
 - Rebekah, Raven, Daniel, Sneha, Kai, Arpan, and David
 - Current enrollment defends at least one more CA hire, and possibly two.
- All CA's have either taken course, previously CA'ed it, or both

CS110 Class Resources I

▪ Course Web Site: <http://cs110.stanford.edu>

- Very simple, optimized to surface exactly what you need and nothing else
- Check the website for information about upcoming lectures, assignment handouts, lecture summaries, and links to lecture slides like the one you're working through right now

▪ Online Student Support

- Peer-collaborative forums: <http://www.piazza.com>
- Staff-managed email queue: cs110@cs.stanford.edu
- My email: jerry@cs.stanford.edu

▪ Office Hours

- My office hours are Tuesdays and Thursdays from 1:00pm until 4:00pm, and they'll normally be held in my [Gates 192](#) office
 - Walk-ins are almost always fine, and I'm happy to help
 - MWF before lecture is not good. I'm a just-in-time kinda lad and I'm almost certainly prepping for class
- CA's will provide a full matrix of office hours, soon to be determined
- Office hours are not for debugging your assignments, and the CAs have been instructed to not look at code. Ever.

CS110 Class Resources II

■ Two Textbooks:

- First textbook is other half of CS107 textbook
 - ["Computer Systems: A Programmer's Perspective"](#), by Bryant and O'Hallaron
 - Stanford Bookstore stocks custom version of just the four chapters needed for CS110
 - Examples in book are in C, though we'll migrate to C++. (Good to know pros and cons of both)
- Second textbook is more about systems-in-the-large, less about implementation details
 - ["Principles of Computer System Design: An Introduction"](#), by by Jerome H. Saltzer and M. Frans Kaashoek
 - Provided free-of-charge online, chapter by chapter
 - Not stocked at Stanford Bookstore by design, since free is better than \$60. You can buy a copy of it from Amazon if you want one

■ Lecture Examples

- Lectures are generally driven by coding examples, and all coding examples can be copied/cloned into local space so you can play and prove they work properly
- Code examples will be developed and tested on the myth machines, which is where you'll complete all of your CS110 assignments
- The accumulation of all lecture examples will be housed in a mercurial repository at **`/usr/class/cs110/lecture-examples`**, which you can initially **`hg clone`**, and then subsequently **`hg pull && hg update`** to get the newer and updated examples as I check them in

■ Lecture Slides

- Will rely on slides when I need to press through lots of information not driven by coding examples
- Most (but not all) lectures will have them. When provided, they'll be organic, in that I'll inject updates and clarifications (and be clear that I added stuff when it really impacts you)
- (Thanks go out to David Mazières, whose excellent CS240H slides inspired my decision to use [markdown](#) and [pandoc](#), and further inspired my decision to, with attribution, steal his stylesheets)

Course Syllabus

■ Overview of Linux Filesystems

- Linux and C libraries for file manipulation: **stat**, **struct stat**, **open**, **close**, **read**, **write**, **readdir**, **struct dirent**, file descriptors, regular files, directories, soft and hard links, programmatic manipulation of them, implementation of **ls**, **cp**, **find**, etc.
- Naming, abstraction and layering concepts in systems as a means for managing complexity, blocks, inodes, inode pointer structure, inode as abstraction over blocks, direct blocks, indirect blocks, doubly indirect blocks, design and implementation of a file system
- Additional systems examples that rely on naming, abstraction, modularity, and layering, including databases, DNS, TCP/IP, network packets, HTTP, REST, descriptors and pids
- Building modular systems with simultaneous goals of simplicity of implementation, fault tolerance, and flexibility of interactions

■ Exceptional Control Flow

- Introduction to multiprocessing, **fork**, **waitpid**, **execvp**, process ids, interprocess communication, context switches, user versus supervisor mode, system calls and how their calling convention differs from those of normal functions
- Protected address spaces, virtual memory, main memory as cache, virtual to physical address mapping, scheduling
- Concurrency versus parallelism, multiple cores versus multiple processors, concurrency issues with multiprocessing, signal masks
- Interrupts, faults, systems calls, signals, design and implementation of a simple shell
- Virtualization as a general systems principle, with a discussion of processes, RAID, load balancers, AFS servers and clients

Course Syllabus (continued)

■ Software-Level Caching

- Expense of system calls, disk seeks, recomputation of $\omega(1)$ algorithms, in-software caching, MRU and LRU techniques, review of profiling tools taught in CS107 and introduction to other ones
- Caching, performance, and consistency as general systems principles, with a discussion of proxy caches, [SPDY](#), [memcached](#), performant web applications (e.g. Facebook) that rely on multiple caching layers, virtual machines

■ C++ STL, C++11 Overview [mostly taught via slides, minimal lecture coverage, possible review session outside of normal lecture time]

- Transition to C++, because it's C and then some, and provides better string facilities, typesafe but flexible containers
- Motivate transition from C to C++, STL **vector**, **list**, **set**, **map**, **unordered_map**, functors, STL hashing, iterators, STL algorithms, new **for** loop model in C++11
- C++11 support for blocks, closures, capture clauses, type inference via **auto**, pass-by-reference, preventing pass-by-value

Course Syllabus (continued)

▪ Threading and Concurrency

- Sequential programming, VLIW concept, desire to emulate the real world within a single process using parallel threads, free-of-charge exploitation of multiple cores (two per **myth** machine, eight per **corn** machine, 24 per **barley** machine), pros and cons of threading versus forking
- C++ threads, **thread** construction using function pointers, blocks, functors, **join**, **detach**, race conditions, **mutex**, IA32 implementation of **lock** and **unlock**, spinlock, busy waiting, preemptive versus cooperative multithreading, **yield**, **sleep_for**
- Condition variables, rendezvous and thread communication, **unique_lock**, **wait**, **notify_one**, **notify_all**, deadlock, thread starvation
- Semaphore concept and **class semaphore** implementation, generalized counters, pros and cons of **semaphore** versus exposed **condition_variable_any**, thread pools, cost of threads versus processes
- Active threads, blocked threads, ready threads, high-level implementation details of a thread manager, **mutex**, and **condition_variable_any**
- Pure C alternatives via **pthread**s, pros and cons of **pthread**s versus C++11's thread package

Course Syllabus (continued)

■ Introduction to Networking

- Client-server model, peer-to-peer model, telnet, protocol as contract for clear communication between programs, request, response, stateless versus keep-alive connections, latency and throughput issues, **gethostbyname**, **gethostbyaddr**, IPv4 versus IPv6, struct sockaddr hierarchy of records, network-byte order
- Ports, sockets, socket descriptors, **socket**, **connect**, **bind**, **accept**, **read**, **write**, simple echo server, time server, concurrency issues, spawning threads to isolate and manage single conversation
- C++ layer over raw C I/O file descriptors, pros and cons, introduction to **sockbuf** and **sockstream** C++ classes (via **socket++** open source project)
- HTTP 1.0 and 1.1, header fields, GET, HEAD, POST, complete versus chunked payloads, response codes, caching
- IMAP protocol, custom protocols, Dropbox and iCloud reliance on HTTP

Course Syllabus (continued)

▪ Additional Topics [as time permits]

- [MapReduce](#) programming model, implementation strategies using multiple threads and multiprocessing
- Nonblocking I/O, where normally slow system calls like **accept**, **read**, and **write** return immediately instead of blocking, **select**, **epoll_*** set of functions, **libev** and **libuv** open source libraries.
- Cross-language development, systems coding in Python, Java, profiling to identify bottlenecks, re-implementing in C or C++, calling from Python, Java.
- Virtualization, revisit virtual memory, threads as virtual processors, virtual file systems, [AGZ]FS, FUSE, virtual runtimes ala JRE and JVM, hardware virtualization ala VMWare.
- Case studies
 - XWindows (legacy, but interesting example of client-server model),
 - Google's MapReduce and Apache's Hadoop,
 - Andrew File System (AFS) and AFS clients,
 - Facebook FBML (implementation in PHP with bridge to Mozilla's C parser),
 - FriendFeed Tornado (implementation in Python with bridge to the Linux **epoll** library)

Student Expectations

■ Programming Assignments

- 45% of final grade
- Expect six or so assignments (depends on how much as-time-permits material I get to)
- Some assignments are single file, others are significant code bases to which you'll contribute. If CS107 is about mastering the periodic table and understanding the chemistry of every single element, CS110 is about building rich, durable polymers
- Lateness discouraged, but forgivable
 - You can grant yourself 24-hour ("late day") extensions, never ever more than two per assignment, Saturdays and Sundays count
 - First three late days are free, each additional one costs you approximately 10%
 - Don't ask for more free late days unless you have an exquisite reason, and the first three were used for equally exquisite reasons.
 - Try to save your free late days because you're actually sick. Those are what your free late days are really for!

Student Expectations

■ Midterm

- Two-hour midterm is Thursday, February 12th at 7:00pm
 - 20% of final grade, material drawn from first five weeks of lecture, mix of implementation and short answer questions
 - Closed-book, closed-notes, closed-electronics, one double-sided cheat sheet that you can prepare ahead of time
 - You must pass the midterm in order to pass the class
 - Passing score will be revealed on midterm solution set, which will be posted well before the withdrawal deadline
 - Practice midterm will be provided a week ahead of time
 - Those with reasonable conflicts can take the exam earlier in the day, provided the two hours you have begin at or after 9:00am and end before or by 5:00pm

■ Final Exam

- Three-hour final is Thursday, March 19th at 8:30am
 - 35% of final grade, cumulative, mix of implementation and short answer questions
 - Closed-book, closed-notes, closed-electronics, two double-sided cheat sheets that you can prepare ahead of time
 - You must pass the final in order to pass the class
 - Multiple practice finals will be provided a week ahead of time
 - No alternate exams
 - Another class via SCPD that meets MWF at 10:00am? RSVP no to their final exam. ☺
 - Another class offered at some other time decided their final is the same time as mine? RSVP no to their final exam. ☺

Honor Code

- Please take it seriously, because the CS Department does.
- Cite all sources and collaborations, and it's much harder to get in trouble.
- The following are clear no-no's
 - Looking at another student's code
 - Showing another student your code
 - Discussing assignments in such detail that you duplicate a portion of someone else's code in your own program
 - Uploading your code to a public repository (e.g. [github](#) or [bitbucket](#)) so that others can easily discover it via word of mouth or search engines.
If you'd like to upload your code to a private repository, you can do so on [bitbucket](#) or some other hosting service that provides free-of-charge private hosting.

Reading

- Skim [Chapter 2, Sections 1 - 4](#)
- Be prepared to consult [Chapter 2, Section 5](#) when Assignment 1 goes live, as it provides the theory that backs what you'll be implementing.
- Live off campus? Read [this](#) so you still can get access to the free version of the online textbook.