

# Announcements

## ▪ Still riding this ECF train!

- I still need to demonstrate and implement the second version of our shell—this one called **simplesh**—to illustrate how a real shell parses the command line, leverages **fork** and **execvp** to spawn off new executables, and exactly what it means for a process to run in the background as opposed to the foreground.
- I also need to work through the client program using and the implementation of the **subprocess**, which demonstrates how different processes can communicate via **pipes**. [#lovethisexample](#)
- I also need to work through the examples from this past [Wednesday's slide deck](#). We haven't gotten to that slide deck yet. [#sadface](#)
- And I need to sneak in an example where **processes** signal other processes. So much to do!
  - Most of our examples just rely on the kernel firing signals at processes that need to be informed about some (possibly bad, possibly not bad) event.
  - However, there are certain times when a process (e.g. your Assignment 3 **tsh** shell, for instance, which you'll get next Friday) needs to signal a child process (e.g. one of the shell's jobs) that something happened. In particular, your **tsh** shell will be the first to hear about **SIGINT** and **SIGTSTP** signals when the user types ctrl-c and ctrl-z. Your **tsh** implementation will want to handle each of those signals by effectively forwarding the same signal to the foreground process group.
- Come later next week, I'll present a straightforward, high-level view of how two (or 200) processes can seemingly run in the same address space by talking about virtual-to-physical memory mapping (our first real foray into virtualization).

# Interprocess communications

- Processes can message other processes using signals via the **kill** system call.

- Prototype:

```
int kill(pid_t pid, int signum);
```

- Analogous to the **/bin/kill** shell command.
  - Unfortunately named, since kill implies **SIGKILL**, which implies death.
  - So named because the default action of most signals in early UNIX implementations was to just terminate the target process (identified by **pid**).
- Returns -1 if the call fails (generally because the calling process doesn't have permission to fire signals at target process), 0 if all is swell.
- **pid** parameter is overloaded to provide many signalling strategies:
  - When **pid** is a positive number, the target is a process with that **pid**.
  - When **pid** is a negative number less than -1, the targets are all processes within the process group **abs(pid)**.
  - **pid** can also be 0 or -1, but we don't need to worry about those (though it's documented in the B&O textbook if you're curious).
- Your Assignment 3 **tsh** shell will use **kill** within your **SIGINT** and **SIGTSTP** handlers to forward those same signals to the foreground process group.

# One final example

## ▪ One final puzzle to ensure you follow the ideas:

- Meaningful example illustrating **kill** function would be too large.
- Instead, we'll work through a small puzzle to confirm you all understand the workflow of the processes and understand how **kill** triggers various signal handlers to be executed. (Error checking is omitted from this example, since it's such a small example, and we'll assume, for simplicity, that nothing ever fails).
- Putting it all together. What're the possible outputs (plural!) of the [following program](#)?

```
static pid_t pid;
static int counter = 0;

static void parentHandler(int unused) {
    counter++;
    printf("counter = %d\n", counter);
    kill(pid, SIGUSR1);
}

static void childHandler(int unused) {
    counter += 3;
    printf("counter = %d\n", counter);
}

int main(int argc, char *argv[]) {
    signal(SIGUSR1, parentHandler);
    if ((pid = fork()) == 0) {
        signal(SIGUSR1, childHandler);
        kill(getppid(), SIGUSR1);
        return 0;
    }

    waitpid(pid, NULL, 0);
    counter += 7;
    printf("counter = %d\n", counter);
    return 0;
}
```

# One final example (continued)

## ■ What's the expected output?

- Make sure you understand why the program on the previous slide is capable of producing two different outputs, depending on processor scheduling.  
(Hint: the child process may or may not exit before **parentHandler** has a chance to kill it!)

```
myth22> ./kill-puzzle  
counter = 1  
counter = 8  
myth22> ./kill-puzzle  
counter = 1  
counter = 3  
counter = 8  
myth22>
```