

CS110 Course Outline

Overview of Linux Filesystems

- Linux and C libraries for file manipulation: `stat`, `struct stat`, `open`, `close`, `read`, `write`, `readdir`, `struct dirent`, file descriptors, regular files, directories, soft and hard links, programmatic manipulation of them, implementation of `ls`, `cp`, `cat`, etc.
- naming, abstraction and layering concepts in systems as a means for managing complexity, blocks, `inodes`, `inode` pointer structure, `inode` as abstraction over blocks, direct blocks, indirect blocks, doubly indirect blocks, design and implementation of a file system.
- additional systems examples that rely on naming, abstraction, modularity, and layering, including DNS, TCP/IP, network packets, HTTP, REST, descriptors and `pids`.
- building modular systems with simultaneous goals of simplicity of implementation, fault tolerance, and flexibility of interactions.

Exceptional Control Flow

- introduction to multiprocessing, `fork`, `waitpid`, `execvp`, process ids, inter-process communication, context switches, user versus supervisor mode.
- protected address spaces, virtual memory, main memory as cache, virtual to physical address mapping.
- concurrency versus parallelism, multiple cores versus multiple processors, concurrency issues with multiprocessing.
- interrupts, faults, systems calls, signals, design and implementation of a simple shell.
- virtualization as a general systems principle, with a discussion of processes, RAID, load balancers, AFS servers and clients.

Software-Level Caching

- expense of system calls, disk seeks, recomputation of $O(1)$ algorithms, in-software caching, MRU and LRU techniques.
- caching, performance, and consistency as general systems principles, with a discussion of proxies, SPDY, `memcached`, performant web applications that rely on multiple caching layers, virtual machines.

Threading and Concurrency

- sequential programming, VLIW concept, desire to emulate the real world with parallel threads, free-of-charge exploitation of multiple cores (two per myth machine, eight per corn machine, 24 per barley machine), pros and cons of `threading` versus `forking`.
- C++ `threads`, `thread` construction using function pointers, blocks, functors, `join`, `detach`, race conditions, `mutex`, IA32 implementation of `lock` and `unlock`, `spinlock`, busy waiting, preemptive versus cooperative multithreading, `yield`, `sleep_for`.
- condition variables, rendezvous and thread communication, `unique_lock`, `wait`, `notify_one`, `notify_all`, deadlock.
- semaphore concept and `class semaphore` implementation, generalized counter, pros and cons of `semaphore` versus exposed condition variables, thread pools, cost of threads versus processes.
- active threads, blocked threads, ready thread queue, high-level implementation details of the thread manager, `mutex`, `condition_variable`, and `condition_variable_any`.
- pure C alternatives via `pthreads`, pros of `pthreads` over C++11 thread package.

Introduction to Networking

- client-server model, peer to peer model, protocol as contract and permitted conversation, request and response as a way to organize modules and their interactions to support a clear set of responsibilities.
- stateless versus keep-alive connections, latency and throughput issues, `gethostbyname`, `gethostbyaddr`, IPv4 versus IPv6, `struct sockaddr` hierarchy of `structs`, network-byte order.
- ports, socket file descriptors, `socket`, `connect`, `bind`, `accept`, `read`, `write`, simple echo server, time server, concurrency issues, spawning threads to isolate and manage single conversation
- C++ layer over raw I/O file descriptors, pros and cons, introduction to `sockbuf` and `sockstream` C++ classes.
- HTTP 1.0 and 1.1, header fields, `GET`, `HEAD`, `POST`, complete versus chunked payloads, response codes, web caching and consistency protocols.
- IMAP, custom protocols, Dropbox and iCloud reliance on variation of HTTP

Additional Topics

- MapReduce programming model, implementation strategies using multiple threads and multiprocessing, comparison to previous systems that do the same thing, but not as well.
- Non-blocking I/O, where normally slow system calls like `open`, `accept`, `read`, and `write` return immediately instead of blocking, `select`, `epoll_*` set of functions, `libev` and `libuv` open source libraries.
- virtualization, briefly revisit virtual memory, threads as virtual processors, virtual file systems, [AGZ]FS, FUSE, virtual runtimes ala JRE and JVM, hardware virtualization ala VMWare.
- cross language development, systems coding in Python, Java, profiling to identify bottlenecks, re-implementing in C or C++, calling from Python, Java.
- case studies: XWindows (legacy, but interesting example of client-server model that goes beyond HTTP), AFS and AFS clients, Facebook FBML (written in PHP with bridge to Mozilla C parser), FriendFeed Tornado (written in Python with bridge to the Linux `epoll` library)

Phil Levis contributed to this handout.