

More Synchronization

CS439: Principles of Computer Systems

February 11, 2015

Last Time

- Hardware Support for Mutual Exclusion and Synchronization
 - read-modify-write instructions
- Mutual Exclusion and Synchronization in Software
 - Abstractions built on top of hardware support
 - Locks
 - Semaphores

Semaphore Recap

- Two types: binary and counted
- Used for locking and synchronization
- Two *atomic* operations:
 - down(): wait for semaphore to be available, decrement value
 - up(): signal waiting threads semaphore is available, increment value
- Need separate semaphores for locking and synchronizing

Semaphore Operations

Semaphore::Down(){

- *if value is == 0, sleep
- *on signal, wake up
- *decrement semaphore value
- *return

}

Semaphore::Up{

- *increment semaphore value
- *if any threads are sleeping on semaphore, wake one of them up
- *return

}

Today's Agenda

- More Synchronization in Software
 - Monitors
 - Locks and Condition Variables
 - Bounded Buffer problem
- Readers/Writers problem

Monitors

Introducing Monitors

- Monitors guarantee mutual exclusion
 - only one thread may execute a given monitor method at a time
- First introduced as a programming language construct (Mesa, Java)
- Now also define a programming convention and can be used in any language (C, C++, ...)
- Object-oriented style:
 - Collect related shared data into an object/module
 - Associate a lock with each one
 - All data is private
 - Define methods for accessing the shared data
 - These are critical sections

Monitors, Formally

A monitor defines a *lock* and zero or more *condition variables* for managing concurrent access to shared data.

- uses the *lock* to ensure that only a single thread is active in the monitor at any point
- the *lock* also provides mutual exclusion for shared data
- *Condition variables* enable threads to block waiting for an event inside of critical sections
 - release the lock at the same time the thread is put to sleep

Monitor Implementation

- Goal is to encapsulate shared data
 - Class in C++ or Java
 - Struct or File in C (logical encapsulation)
- Goal is to provide for mutual exclusion
 - Has a lock (exactly one!)
- Goal is to allow for synchronization
 - Has *condition variables*
- Goal is to allow operations on the shared data
 - Has functions or methods

Monitor Functions: Implementation

- Acquire the lock at the start of every function (first thing!)
 - Operate on the shared data
 - Temporarily release the lock if they can't complete due to missing resource (use condition variable for this)
 - Reacquire the lock when they can continue (again, condition variable!)
 - Operate on the shared data
- Release the lock at the end

Semaphore Example:

Producers/Consumers (Recall)

```
Semaphore mutex = 1    //access to buffer
Semaphore empty = N    //count of empty slots
Semaphore full = 0      //count of full slots
int buffer[N]
```

```
BoundedBuffer::Producer(){
    <produce item>
    empty->down() //get empty spot
    mutex->down() //get access to buffer

    <add item to buffer>

    mutex->up() //release buffer
    full->up() //another item in buffer
}
```

```
BoundedBuffer::Consumer(){
    full->down()    //get item
    mutex->down() //get access to buffer

    <remove item from buffer>

    mutex->up() //release buffer
    empty->up() //another empty slot
    <use item>
}
```

```
Semaphore mutex = 1 //access to buffer
Semaphore empty = N //count of empty slots
Semaphore full = 0 //count of full slots
int buffer[N]
```

```
BoundedBuffer::Producer(){
    <produce item>
    empty->down() //get empty spot
    mutex->down() //get access to buffer

    <add item to buffer>

    mutex->up() //release buffer
    full->up() //another item in buffer
}
```

Condition Variables

- Enable threads to wait efficiently for changes to shared state protected by a lock
- Each one is a queue of waiting threads (no state!)
- Enable the thread to sleep inside a critical section by *atomically* releasing the Lock at the same time the thread is put to sleep

Rule: A thread *must* hold the lock when doing condition variable operations

Condition Variable Operations

1. Wait(Lock lock)

- atomic (release lock, move thread to waiting queue, suspend thread)
- when the thread wakes up it re-acquires the lock (before returning from wait)
- thread will *always* block

2. Signal(Lock lock)

- wake up waiting thread, if one exists. Otherwise, it's a no-op

3. Broadcast(Lock lock)

- wake up all waiting threads, if any exist. Otherwise, it's a no-op

Monitor Operations

Lock->Acquire() //acquires lock, when returns, thread has lock

Lock->Release() //releases lock

CondVar::Wait(lock){

- *move thread to waiting queue
and suspend thread

- *release lock

- *on signal, wake up, re-acquire
lock

- *return

}

CondVar::Signal(){

- *wake up a process waiting on
condVar

- *return

}

CondVar::Broadcast(){

- *wake up ALL processes waiting
on condVar

- *return

}

Resource Variables

- Conditions variables (unlike semaphores) keep no state
- Each condition variable should have a resource variable that tracks the state of that resource
 - *You must maintain this variable*
- Check the resource variable before calling wait on the associated condition variable to ensure the resource really *isn't* available
- Once the resource is available, claim it (subtract the amount you are using!)
- Before signaling that you are finished with a resource, indicate the resource is available by increasing the resource variable

Monitor Example: Items in a Queue

```
Lock lock;  
Condition fullCV;  
int queue_size;  
  
void Add(item){  
    lock->acquire()  
    put item on queue;  
    queue_size++;  
    fullCV->signal(lock)  
    lock->release()  
}
```

```
void remove(){  
    lock->acquire()  
  
    while(queue_size <=0)  
        fullCV->wait(lock);  
    queue_size--;  
  
    remove item;  
    lock->release();  
}
```

Signal() Semantics

Which thread executes once signal() is called?

- If there are no waiting threads, the signaler continues and the signal is effectively lost
- If there is a waiting thread (or two):
 - There are at least two ready threads: the one that called signal() and the one that was (or will be) awakened
 - Exactly one of the threads can execute or we will have more than one thread active in the monitor (violates mutual exclusion!)
 - So which thread gets to execute?

Whose turn is it?

Mesa/Hansen Style

- The thread that signals keeps the lock (and thus the processor)
- The waiting thread waits for the lock
 - Signal is only a hint that the condition may be true: shared state may have changed!
 - Adding signals affects performance, but never safety
- Implemented in Java and most real operating systems

Hoare Style

- The thread that signals gives up the lock and the waiting thread gets the lock
 - Signaling is atomic with the resumption of the waiting thread
 - Shared state cannot change before waiting thread is resumed
- When the thread that was waiting and is now executing exits or waits again, it releases the lock back to the signaling thread
- Implemented in most textbooks (not yours!)

More on Turns

- With Mesa-style, waiting thread may need to wait again after it is wakened (Why?), so *while* is important
- With Hoare-style, we can change *while* to *if* because a waiting thread runs immediately

```
public void remove(){  
    lock->Acquire()  
    while (queue_size <= 0){  
        full->wait(lock);  
    }  
    queue_size--;  
    remove item;  
    lock->Release();  
}
```

Regardless, if you assume there is NO atomicity between `signal()` and the return from `wait()`, your code will always work.

Semaphore Example:

Producers/Consumers (Recall)

```
Semaphore mutex = 1    //access to buffer
Semaphore empty = N    //count of empty slots
Semaphore full = 0      //count of full slots
int buffer[N]
```

```
BoundedBuffer::Producer(){
    <produce item>
    empty->down() //get empty spot
    mutex->down() //get access to buffer

    <add item to buffer>

    mutex->up() //release buffer
    full->up() //another item in buffer
}
```

```
BoundedBuffer::Consumer(){
    full->down()    //get item
    mutex->down() //get access to buffer

    <remove item from buffer>

    mutex->up() //release buffer
    empty->up() //another empty slot
    <use item>
}
```

Semaphore Example:

Producers/Consumers (Recall)

```
Semaphore mutex = 1    //access to buffer
Semaphore empty = N    //count of empty slots
Semaphore full = 0      //count of full slots
int buffer[N]
```

```
BoundedBuffer::Consumer(){
    full->down()    //get item
    mutex->down() //get access to buffer

    <remove item from buffer>

    mutex->up() //release buffer
    empty->up() //another empty slot
    <use item>
}
```

Monitor Example: Producers/Consumers

```
class BBMonitor{
  public: <methods>
  private: item buffer[N];
          Lock lock;
          Condition fullCV, emptyCV;
          int empty=N;
          int full=0;
}

BoundedBuffer::Producer(){
  lock->Acquire()

  while(empty == 0)
    emptyCV->Wait(lock) //get empty spot
  empty--;

  <add item to buffer>

  full += 1
  fullCV->Signal(lock) //another item in buffer
  lock->Release()
}
```

```
BoundedBuffer::Consumer(){
  lock->Acquire()

  while(full == 0)
    fullCV->Wait(lock) //get item
  full--

  <remove item from buffer>

  empty++;
  emptyCV->Signal(lock) //another empty slot
  lock->Release()
}
```

iClicker Question

Every monitor function should begin with what command?

- A. Wait()
- B. Signal()
- C. Lock->Acquire()
- D. Lock->Release()
- E. Broadcast()

Monitor Summary

- A monitor wraps operations with a lock
- Condition variables release lock temporarily
- Monitors can be implemented by following the monitor rules for acquiring and releasing locks

Comparing Monitors and Semaphores

- Condition variables do not have any history
 - on `signal()` if no one is waiting, the signal is a no-op
 - if thread then calls `condition->wait()`, it waits.
- Semaphores do have history
 - on `up()` if no one is waiting, the value of the semaphore is incremented
 - if a thread then calls `semaphore->down()`, the value is decremented and the thread continues

So... signal() and down()

- In semaphores, down() and up() are commutative
 - result is the same regardless of the order of execution
- Condition variables are not commutative
 - so they must be in a critical section to access state variables and do their job

You *can* implement Monitors with Semaphores.

Monitors and Semaphores: Recap

Both

- Provide mutual exclusion and synchronization
- Have signal() and wait()
- Support a queue of processes that are waiting to access a critical section (e.g., to buy milk)
- No busy waiting!

Semaphores

- Semaphores are basically generalized locks
- Binary and counting semaphores
- Used for mutual exclusion and synchronization

Monitors

- Consist of a lock and one or more condition variables
- Encapsulate shared data
- Use locks for mutual exclusion
- Use condition variables for synchronization
- Wrap operations on shared data with a lock
- Condition variables release lock temporarily

A Different Type of Problem

- We've looked at problems where we protect shared data by only allowing one thread in the critical section at a time
- Is this always appropriate? When might we want to let more threads access shared data at once?

Readers/Writers Problem

- Data is shared among several threads
 - Some only read
 - Some only write
- To get correct results, we allow *multiple readers* at a time, but only *one writer* at a time
- How can we control access to the object to permit this protocol?

Correctness Criteria

- Each read or write of the shared data *must* happen within a critical section
- Guarantee mutual exclusion for writers
- Allow multiple readers to execute in the critical section at once
- Allow one writer (and no readers) to execute in the critical section at once

Readers and Writers: Monitor Solution

- What methods do we need?
 - How many locks?
 - How many condition variables?
 - What should we name them?
 - Any other variables?
-
- Assume we're going to say <read> and <write> for accesses to the shared data.

Readers and Writers: Monitor Solution

Variables:

write(){

read(){

Is our solution fair?

}

}

A. Yes

B. No, favors readers

C. No, favors writers

Understanding Our Solution

It works, but it favors readers over writers

- Any reader blocks all writers
- All readers must finish before a writer can start
- Last reader will wake any writer, but a writer wakes all readers and writers
- If a writer exits and a reader goes next, then all readers that are waiting will get through

Readers and Writers: Monitor Solution

Variables:

write(){

read(){

}

}

Alternative Semantics

- It may be that you would like a writer to enter its critical section as soon as possible.
- How could we implement that?

Signal vs. Broadcast

- It is always safe to use `broadcast()` instead of `signal()`
 - only performance is affected
- `signal()` is preferable when
 - at most one waiting thread can make progress
 - any thread waiting on the condition variable can make progress
- `broadcast()` is preferable when
 - multiple waiting threads may be able to make progress
 - the same condition variable is used for multiple predicates
 - some waiting threads can make progress, others can't

Summary

- A monitor wraps operations with a lock
- Condition variables release lock temporarily
- Monitors do not keep state---a call to wait() will *always* wait
- Monitors can be implemented by following the monitor rules for acquiring and releasing locks

Announcements

- Homework 3 due Friday 9:45a
- Project 1 is posted due 2/27
- Exam 1 is in two weeks!
 - Wednesday, 2/25! 7p! UTC 2.112A!