

Heap Memory Management

CS439: Principles of Computer Systems

March 11, 2015

Last Time

- Paging Mechanism
 - Page Faults
- Paging Policies
 - Replacement algorithms
 - FIFO, Optimal, LRU, Clock, Second Chance
 - Local vs. Global

Today's Agenda

- Paging Policies
 - Load Control Strategies
 - Page Sizes
- Heap Memory Management
 - Explicit vs Automatic/Implicit
 - Allocation techniques
 - Contiguous allocation (bump pointer)
 - Free lists (analogous to pages in memory)
 - Explicit deallocation

Virtual Memory

Thrashing

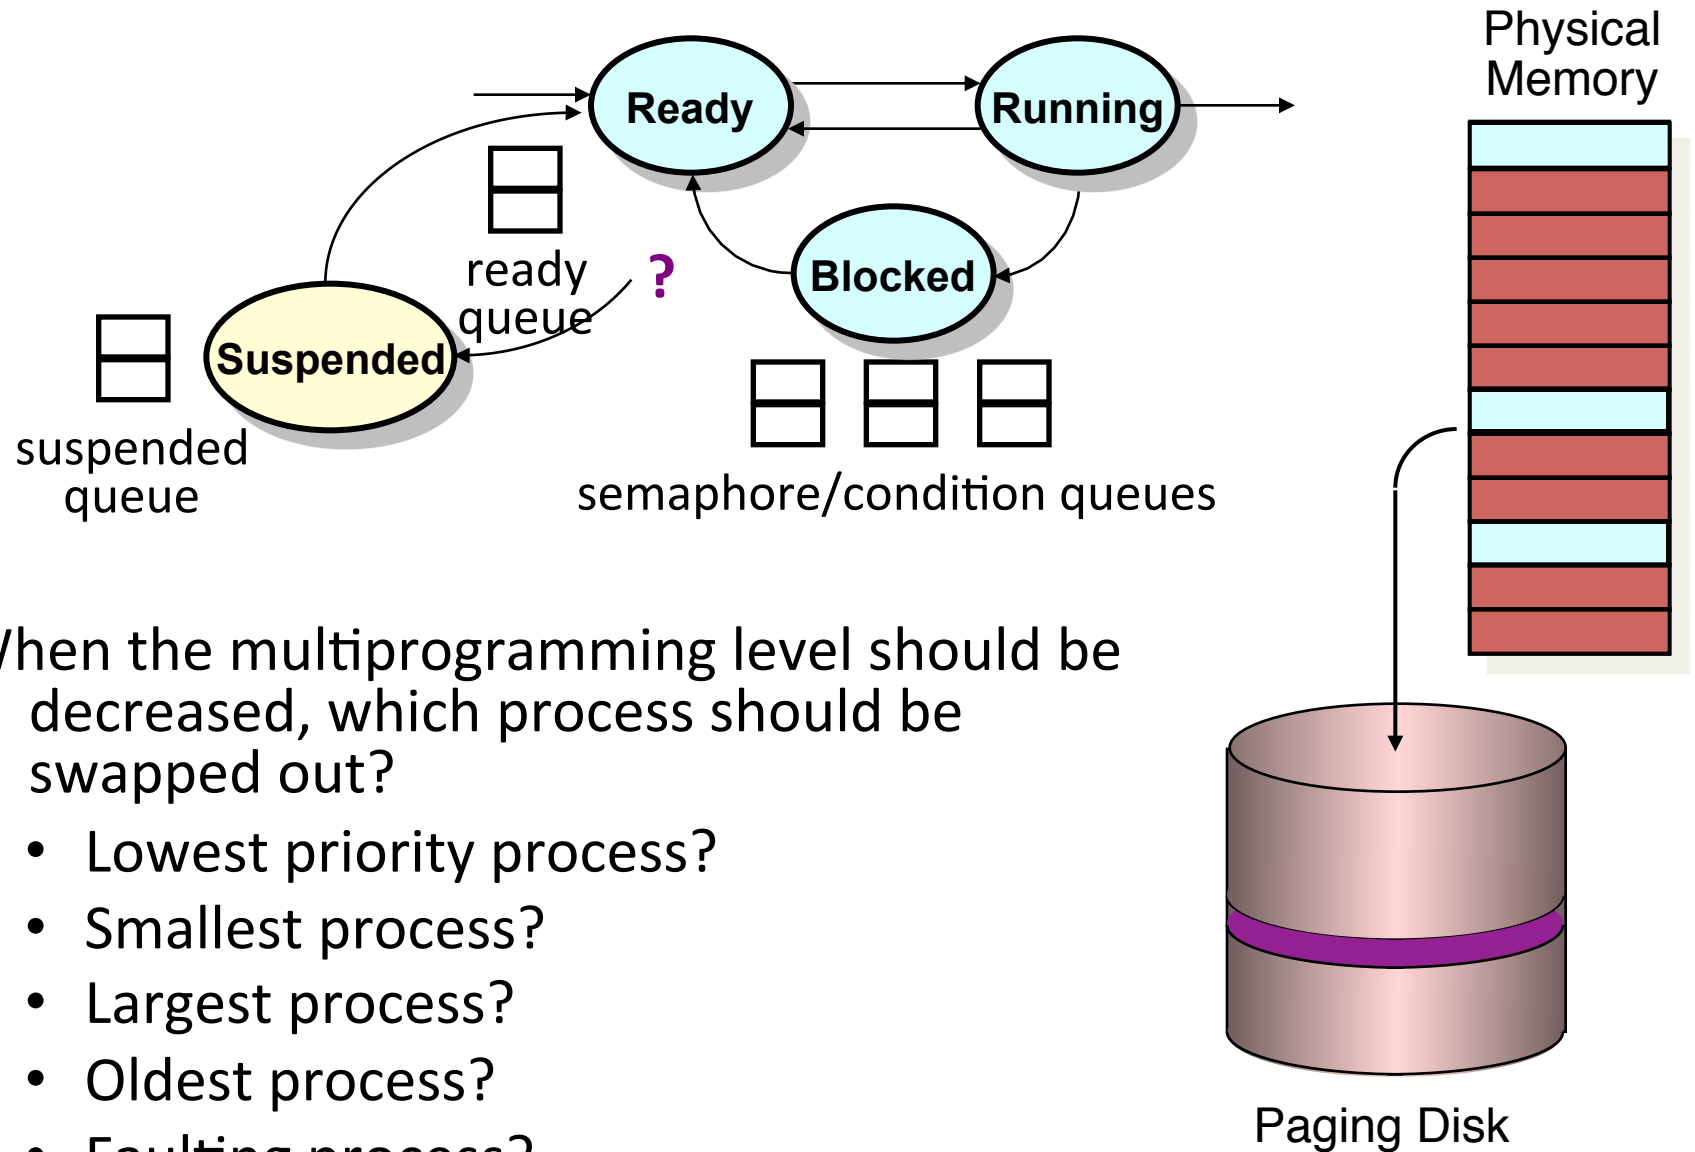
- *Thrashing* occurs when the memory is over-committed and pages are tossed out while they are still in use
- Many memory references cause pages to be faulted in
 - Very serious and very noticeable loss of performance

How do we limit thrashing in a multiprogrammed system?

Load Control

- *Load control* refers to the number of processes that can reside in memory at one time
- Working set model provides implicit load control by only allowing a process to execute if its working set fits in memory
- BUT process frame allocations are variable
- What happens when the total number of pages needed is greater than the number of frames available?
 - Processes are swapped out to disk

Load Control



Load Control: Text Description

- When a process is totally swapped out of memory it is put onto swap (aka the paging disk)
- Adds another stage to the process life cycle
 - This new stage is called suspended
 - We also saw this stage in relocation, when we also swapped out entire processes
 - Can go from any of the other states to suspended
 - Usually blocked or ready
 - Process can go from suspended to ready

Another Decision: Page Sizes

Page sizes are growing slowly but steadily. Why?

- Benefits for small pages: more effective memory use, higher degree of multiprogramming possible
- Benefits for large pages: smaller page tables, reduced I/O time, fewer page faults
- Growing because:
 - memory is cheap---page tables could get huge with small pages and internal fragmentation is less of a concern
 - CPU speed is increasing faster than disk speed, so page faults cause a larger slow down

iClicker Question

Can an application modify its own translation tables (however they are implemented)?

A. Yes

B. No

Summary: Paging

We've considered:

- Placement Strategies
 - None needed, can place pages anywhere
- Replacement Strategies
 - What to do when more jobs exist than can fit in memory
- Load Control Strategies
 - Determine how many jobs can be in memory at one time

Summary:

Paging

The Good

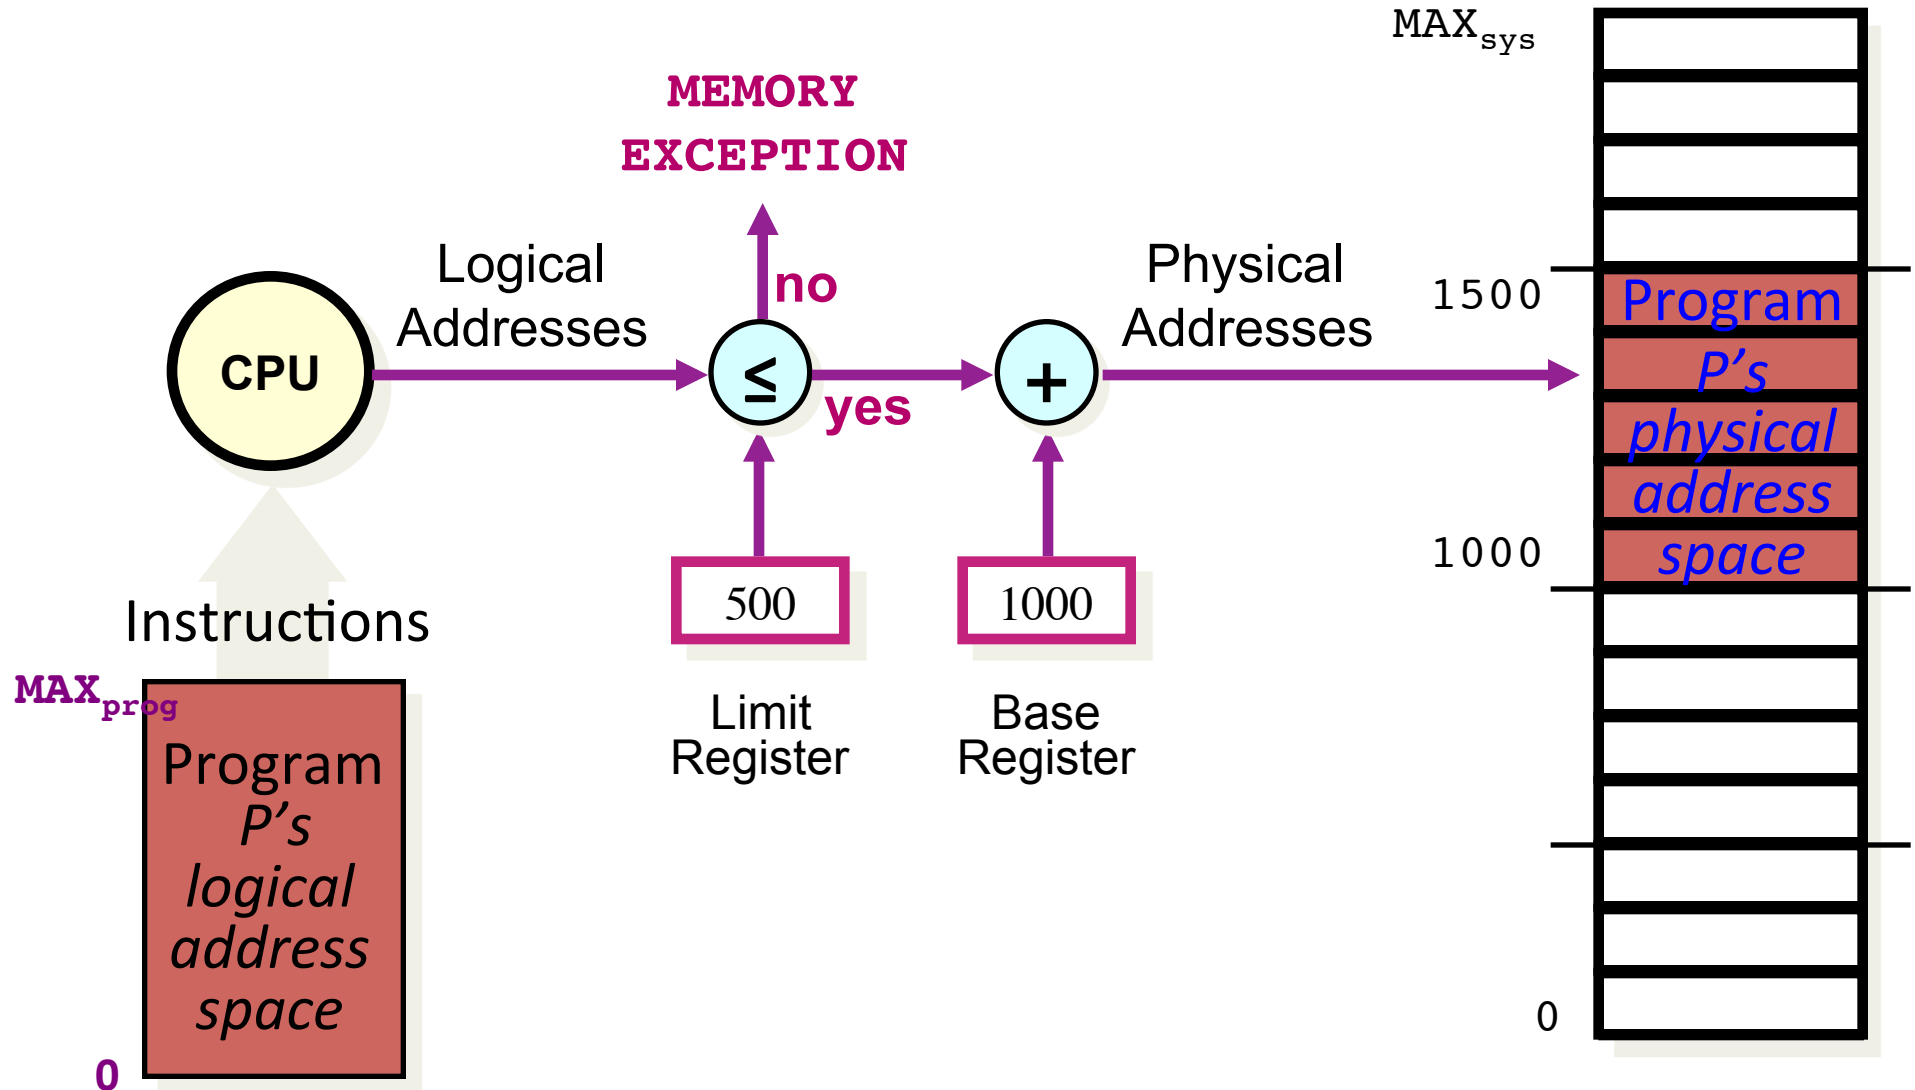
- Eliminates the problem of external fragmentation
- Allows sharing of memory pages amongst processes
- Enables processes to run when they are only partially loaded into main memory

The Cost

- Translating from a virtual address to a physical address is time consuming
- Requires hardware support (TLB) to be decently efficient
- Requires more complex OS to maintain the page table

The expense of memory accesses and the flexibility of paging make paging cost effective.

Dynamic Relocation: Address Translation

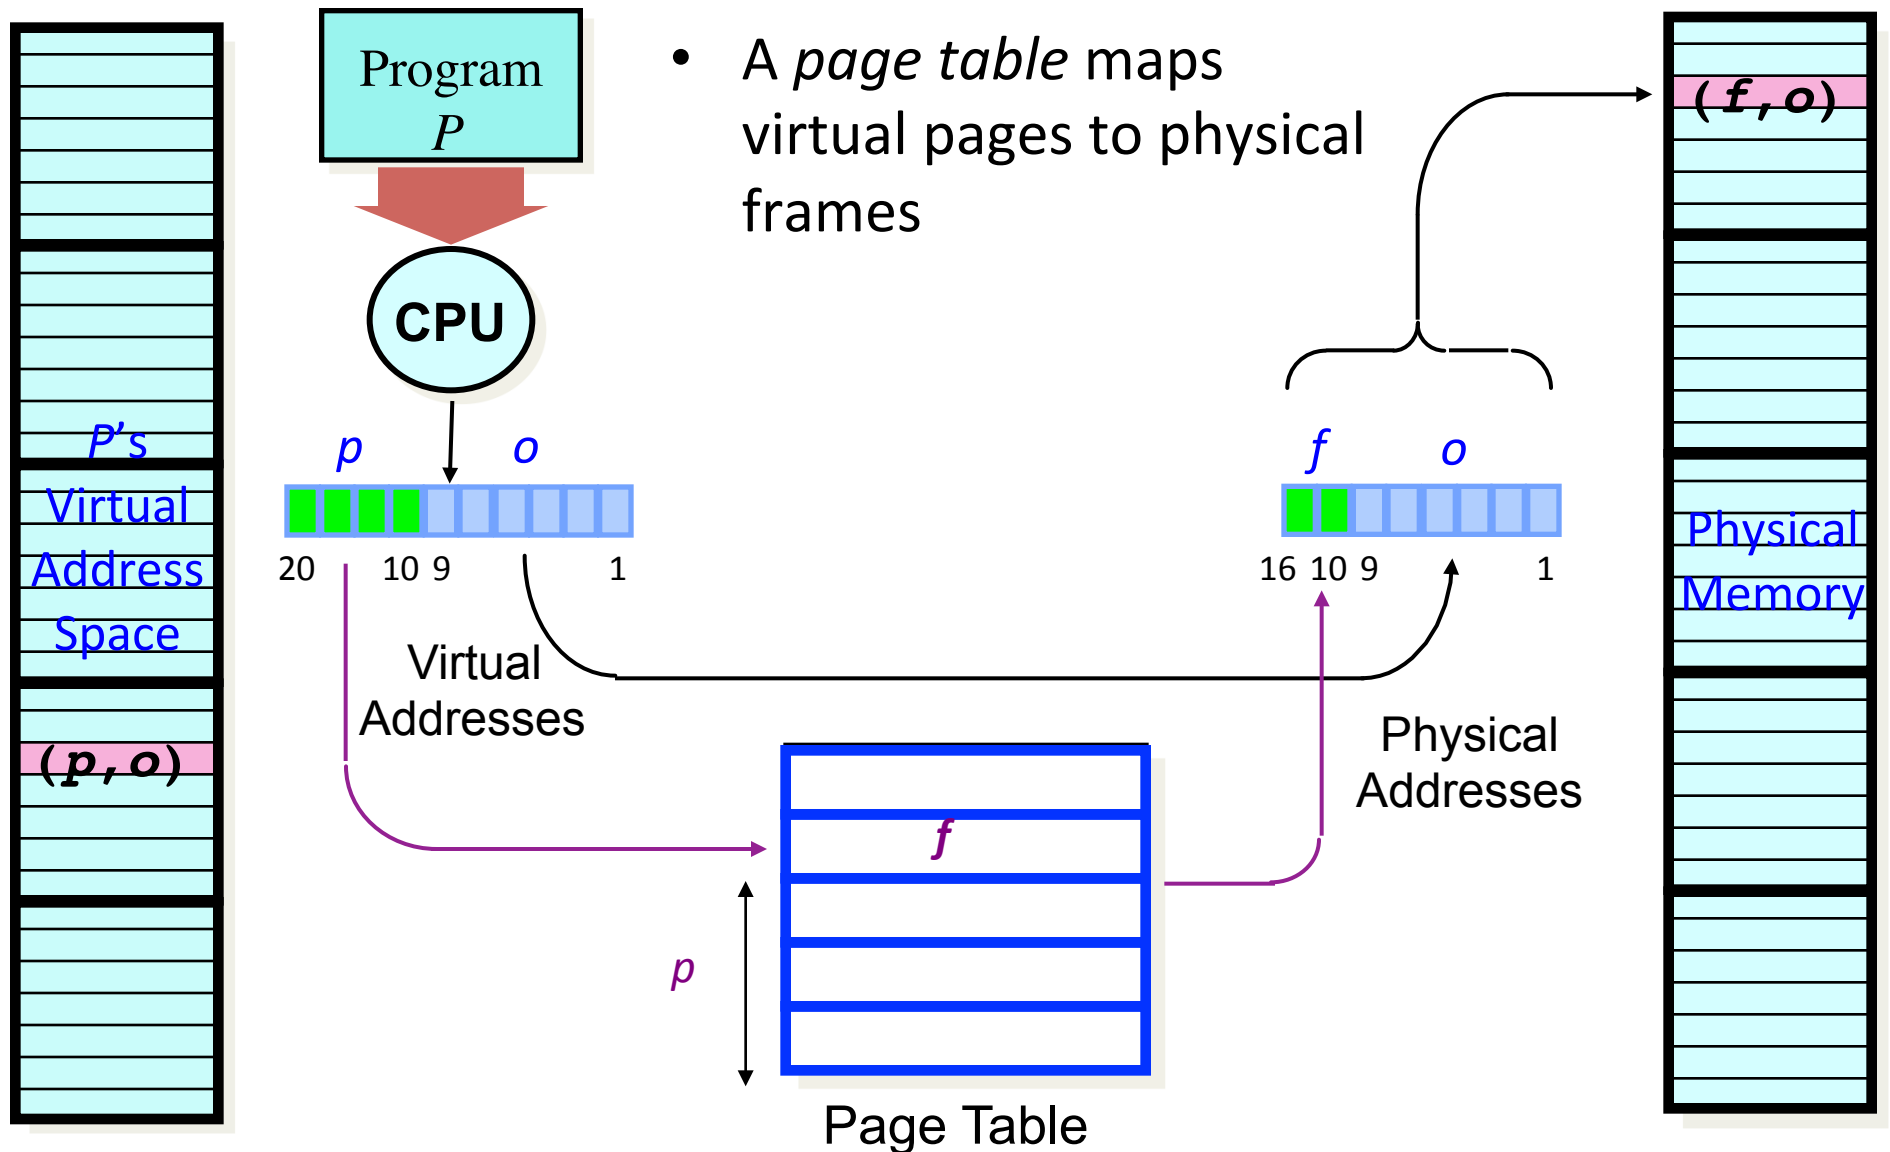


Dynamic Relocation:

Text Description

- How we get from a program's logical address space to physical memory.
- The bound register is used to make sure that the program isn't trying to access memory outside of its space. For example, if the programs logical address space ranges from 0 to 500 then the bound register would hold a value of 500.
- The base register holds the beginning of that program's addresses in physical memory. So if the programs memory started at address 1000 then the base register would hold 1000.
- Steps to Address Translation:
 - The program gives a logical address, or instruction, to the CPU.
 - The MMU now does 2 things at once:
 - It checks the address against the bound register, if the address is greater than the bound register a memory exception is thrown. This exception is a hardware interrupt that will be handled by the OS. This exception indicates the program was trying to access something that doesn't belong to it.
 - It adds the base register to the logical address to get the physical address. For example logical address 8 would become 1008 if the base register was 1000.
- NOTE: Even when the address is successfully translated, you are not protected from memory errors: you could still accidentally try to access memory in your address space that hasn't be initialized yet.

Virtual Address Translation



Memory Management:

Putting it all Together

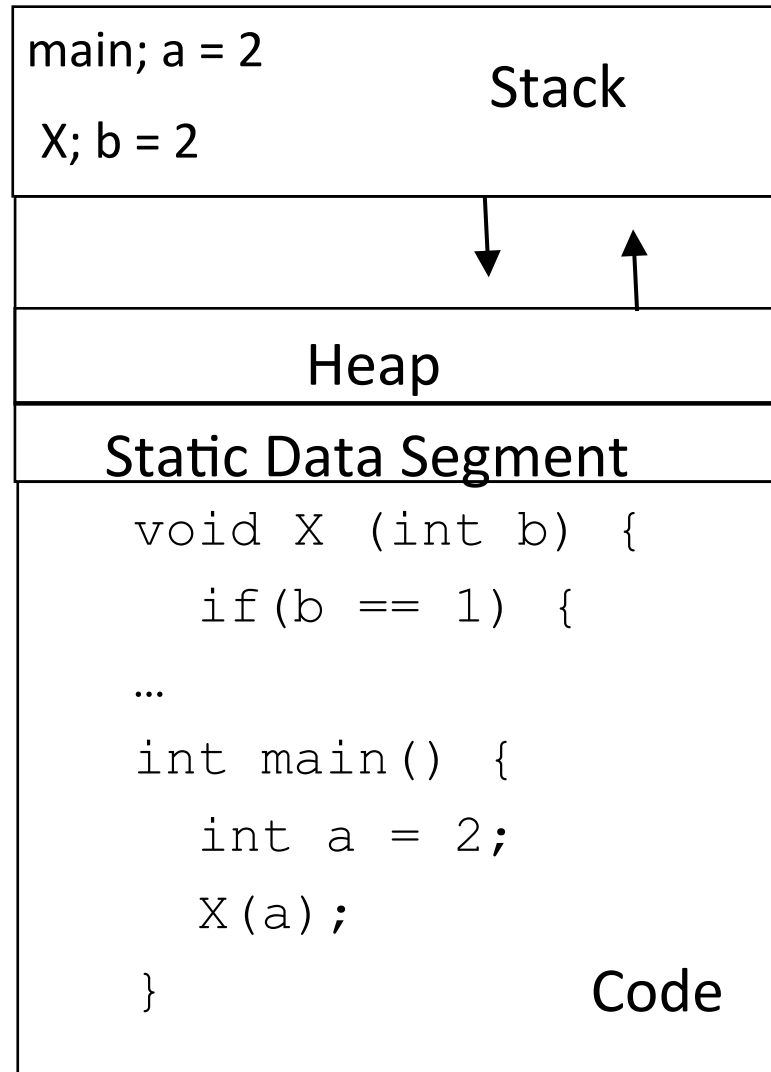
- Dynamic Relocation with Base and Bounds:
 - Simple, but inflexible
 - Degree of multiprogramming limited, memory limited to physical memory size, no sharing of memory, memory allocation/deallocation difficult
 - Use compaction to solve external fragmentation
- Paging:
 - Process generates virtual addresses from 0 to Max
 - OS divides processes into pages
 - manages a page table for every process
 - manages the pages in memory
 - Simplifies memory allocation since any page can be allocated to any frame
 - Page tables can be very large
 - Page Replacement Algorithms
 - FIFO, Optimal, LRU, Clock, Enhanced Clock, Working Set
 - Design Considerations (page size, global vs. local, ...)

Heap Memory Management

Heap Memory Management

- Where and how do we manage dynamically allocated (user) memory (a.k.a. the heap!)?
 - Program/runtime system requests memory from the OS for the heap
 - OS gives memory to a process 1 to k pages at a time (Why?)
 - The runtime system manages the heap memory
 - Typically, the memory is not returned to the OS until the program ends
- How does the runtime system *efficiently* create and recycle memory on behalf of the program?
 - What makes this problem important?
 - What makes this problem hard?

Reminder



What's in the heap?

- Dynamically allocated program objects and data
- Needed when required memory size is not known until the program runs

Two Categories of Heap Memory Management

- Explicit memory management
 - The program(mer) explicitly manages all of the memory
 - Allocation: malloc/new
 - Deallocation: free/delete
 - Pointers: anything may or may not be a pointer
 - Example languages: C, C++
- Automatic memory management (Garbage Collection)
 - The program(mer) explicitly allocates memory, but the runtime system manages it
 - Allocation: new
 - Deallocation: *None*
 - Pointers: Program and runtime system know all pointers
 - Example languages: Java, ML, Python

Key Issues

- How to allocate the memory
 - How to organize the memory space
 - Fast allocation
 - Low fragmentation (wasted space)
- How to deallocate the memory
 - Fast reclamation
 - Discriminating live (in use) objects and garbage (automatic memory management only)

Explicit Memory Management

Two Pieces

- User
 - Explicitly allocates memory by requesting a number of bytes
 - May explicitly request deallocation of memory when it is no longer used
- Runtime System
 - Receives requests for memory
 - Identifies appropriate location for allocation
 - If allocation doesn't fit, requests more memory from the Operating System
 - Returns pointer
 - Later, frees allocation on request

Runtime System Requirements

- Handle arbitrary request sequences
 - Memory may be allocated and freed in any order
- Make immediate responses to requests
 - Cannot reorder/buffer requests to improve performance
- Use only the heap
 - Any data structures (such as free list) used by `malloc()`/`free()` must be stored on the heap
- Align blocks (e.g., on 8-byte boundary)
 - Blocks must be able to hold any type of data object
- Not modify allocated blocks
 - Can only manipulate or change free blocks
 - Cannot modify (or move!) other blocks after they are allocated
 - Results in fairly simple allocation policies

Allocation Techniques


1. Bump-pointer

- Contiguous allocation (for all requested blocks)
- Pointer begins at start of heap
- As requested, bytes allocated, and pointer is “bumped” past allocation

Bump Pointer (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```


```
➔ char *p1 = malloc(3);
   char *p2 = malloc(1);
   char *p3 = malloc(4);
   free(p2);
   char *p4 = malloc(6);
   free(p3);
   char *p5 = malloc(2);
   free(p1);
   free(p4);
   free(p5);
```



Bump Pointer (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```



```
➡ char *p1 = malloc(3);
   char *p2 = malloc(1);
   char *p3 = malloc(4);
   free(p2);
   char *p4 = malloc(6);
   free(p3);
   char *p5 = malloc(2);
   free(p1);
   free(p4);
   free(p5);
```



Bump Pointer (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```


```
char *p1 = malloc(3);
char *p2 = malloc(1);
➔ char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Bump Pointer (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```


```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
→ free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Bump Pointer (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```

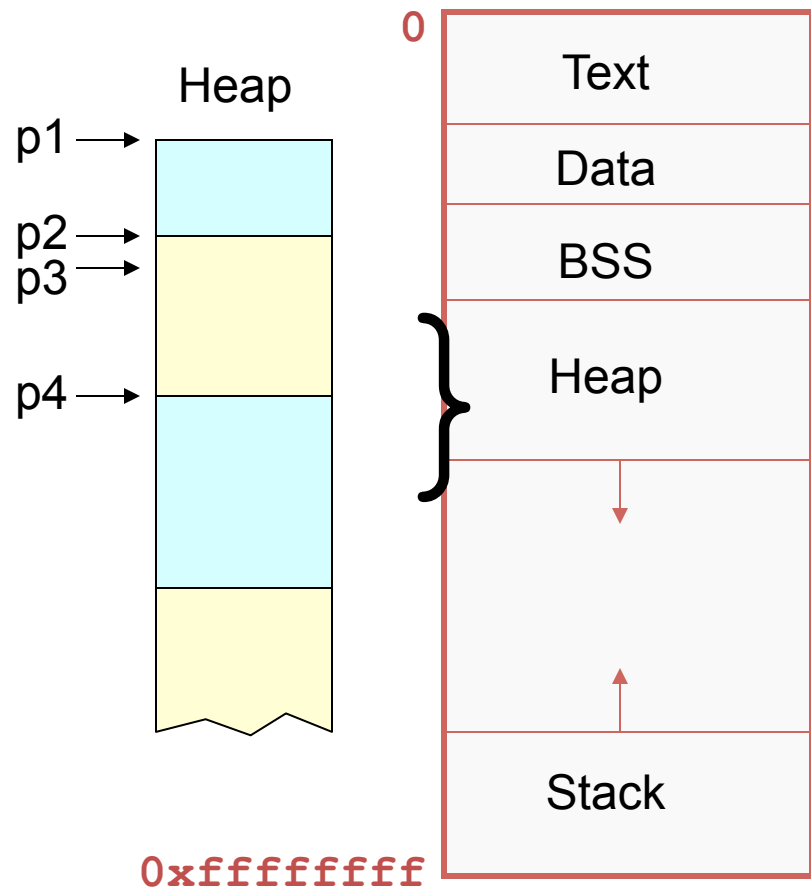
```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
➔ char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Bump Pointer (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```

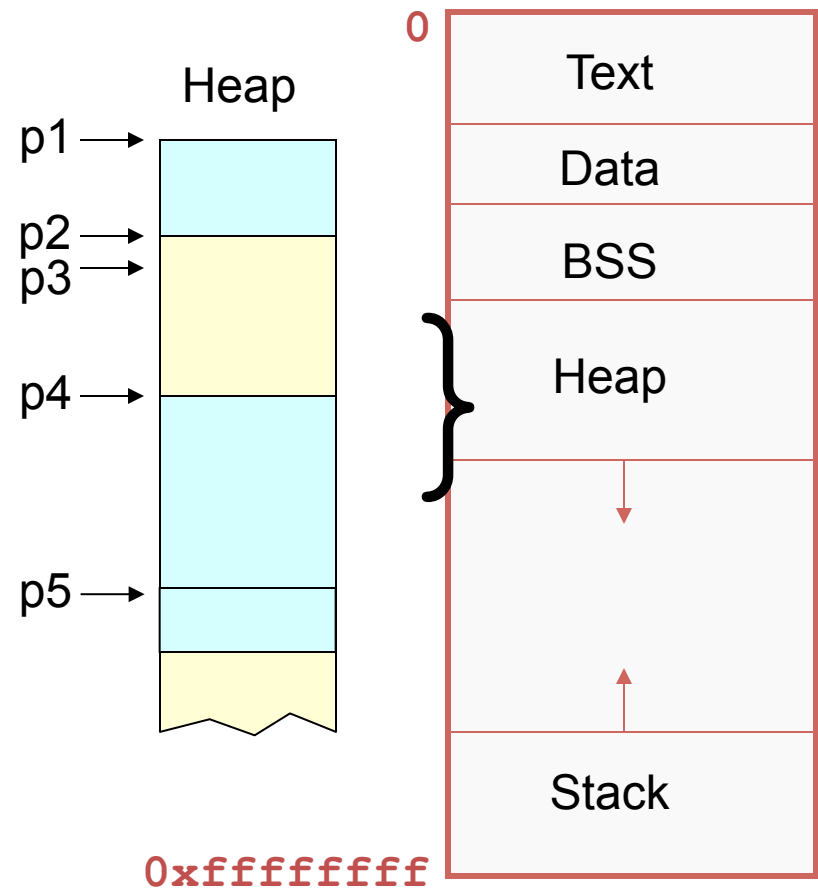
```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
➡ free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Bump Pointer (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```


```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
➔ char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Bump Pointer (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```


```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
→ free(p1);
  free(p4);
  free(p5);
```



Bump Pointer (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```



```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Bump Pointer (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```

```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



0xffffffff

Bump Pointer (C Example): Text Description

- In this address space, the heap grows downwards towards higher addresses
 - order of segments starting at address 0: text, Data, BSS and heap
 - stack starts at some high address later and grows toward lower addresses
 - note: stack and heap grow towards each other
- List of C commands and their consequences:
 - note: originally heap is empty and bump pointer at address 100
 - `char *p1 = malloc(3)`
 - 3 bytes are allocated on the heap
 - p1 holds address 100, the beginning of the 3 bytes
 - bump pointer moved to address 103
 - `char* p2 = malloc(1)`
 - 1 byte is allocated on the heap
 - p2 holds address 103, the beginning of the 1 byte
 - bump pointer moved to address 104
 - `char* p3 = malloc(4)`
 - 4 bytes are allocated on the heap
 - p3 holds address 104, the beginning of the 4 bytes
 - bump pointer moved to address 108
 - `free(p2)`
 - memory that p2 points to is deallocated
 - now is a gap in the heap from 103 to 104
 - gap stays because bump pointer allocation doesn't move memory around once it has been allocated
 - `char* p4 = malloc(6)`
 - 6 bytes allocated on the heap
 - p4 holds address 108, the beginning of the 6 bytes
 - bump pointer moved to address 114
 - `free(p3)`
 - memory that p3 points to is deallocated
 - now is a gap in the heap from 103 to 108
 - this new gap is coalesced with the old gap created by freeing p2
 - `char* p5 = malloc(2)`
 - 2 bytes allocated on the heap
 - p5 points to address 114, the beginning of the 2 bytes
 - bump pointer is moved to address 116
 - `free(p1)`
 - memory that p1 points to is deallocated
 - now a gap in heap from 100 to 108
 - `free(p4)`
 - memory that p4 points to is deallocated
 - now gap in heap from 100 to 114
 - `free(p5)`
 - memory that p5 points to is deallocated
 - now a gap in heap from 100 to 116
 - now: all of the memory that was originally allocated on the heap has been reclaimed
 - note: pointers still hold their old addresses because they were not NULLed out
 - this means can still try to use them but will get a segfault

Allocation Techniques

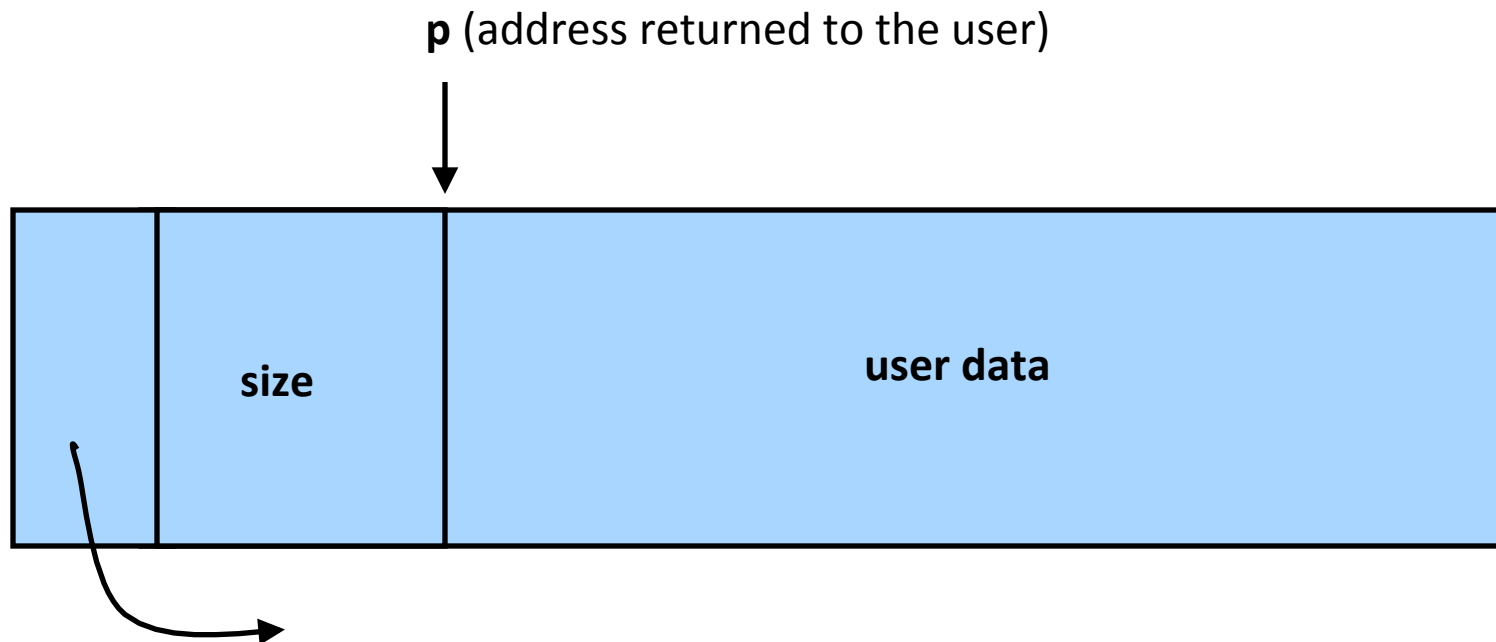
2. Free List

- Divides memory into some size blocks
- Maintains a free list
 - Must be stored in the heap
 - Uses a special structure for free blocks
- To allocate memory, find block in the free list
 - Using what algorithm? Guess!
 - If the right size does not exist, carves up a bigger piece
- To deallocate memory, put memory back on the free list

Free Block: Pointer, Size, Data


Free block in memory

- Pointer to the next free block
- Size of the free block
- Free space (that can be allocated to user)



Free List: Circular Linked List

- Free blocks, linked together
 - Example: circular linked list
- List may be ordered by address or by size, depending on allocation algorithm




Choosing the Spot

First Case: A Perfect Fit

Suppose the block is a perfect fit


- Remove the element from the list
- Link the previous element with the next element
- Return the current element to the user (skipping header)



Choosing the Spot

Second Case: Block is Too Big


- Suppose the block is bigger than requested
 - Divide the free block into two blocks
 - Keep first (now smaller) block in the free list
 - Allocate the second block to the user



Free List (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```

```
➔ char *p1 = malloc(3);
   char *p2 = malloc(1);
   char *p3 = malloc(4);
   free(p2);
   char *p4 = malloc(6);
   free(p3);
   char *p5 = malloc(2);
   free(p1);
   free(p4);
   free(p5);
```




Free List (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```

➔


```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Free List (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```


```
char *p1 = malloc(3);
char *p2 = malloc(1);
➔ char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Free List (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```

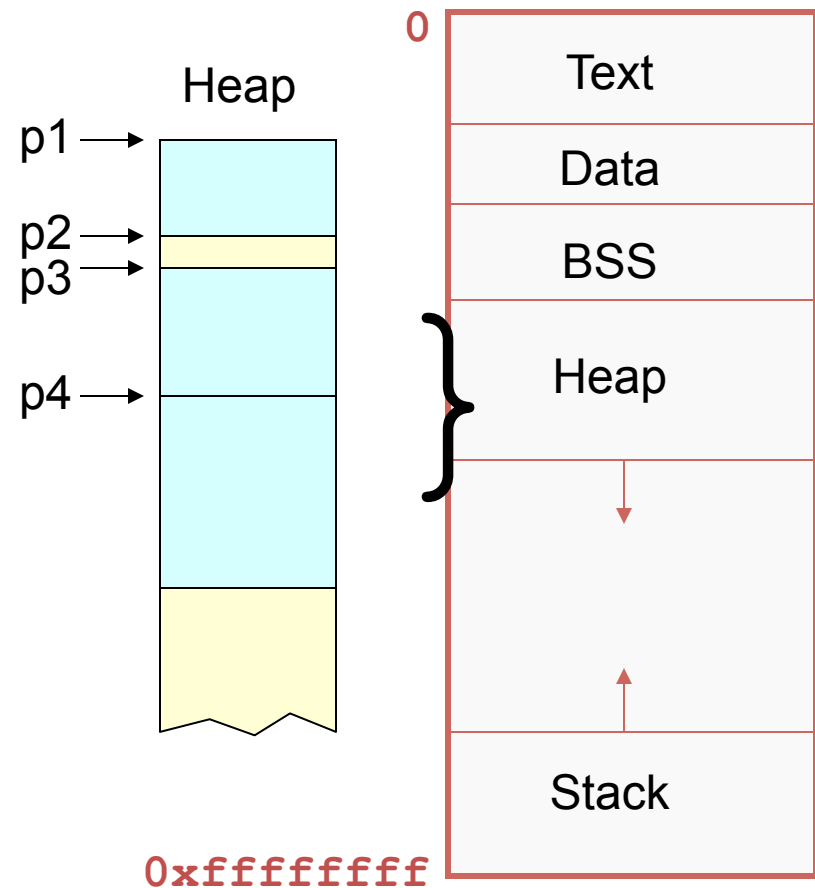
```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
→ free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Free List (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```

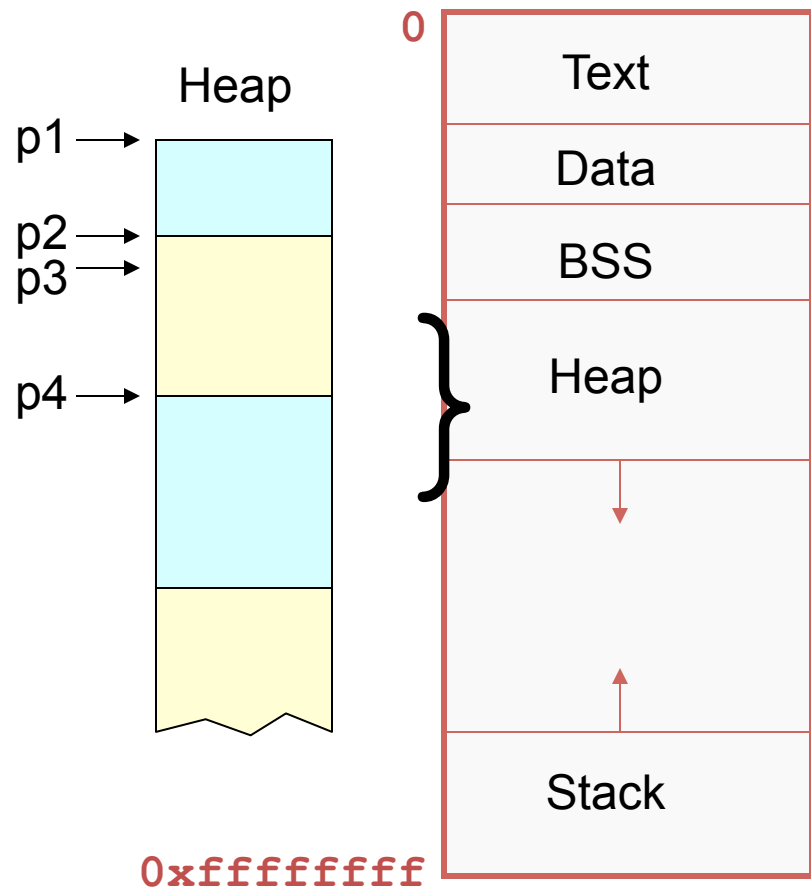
```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
➔ char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Free List (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```

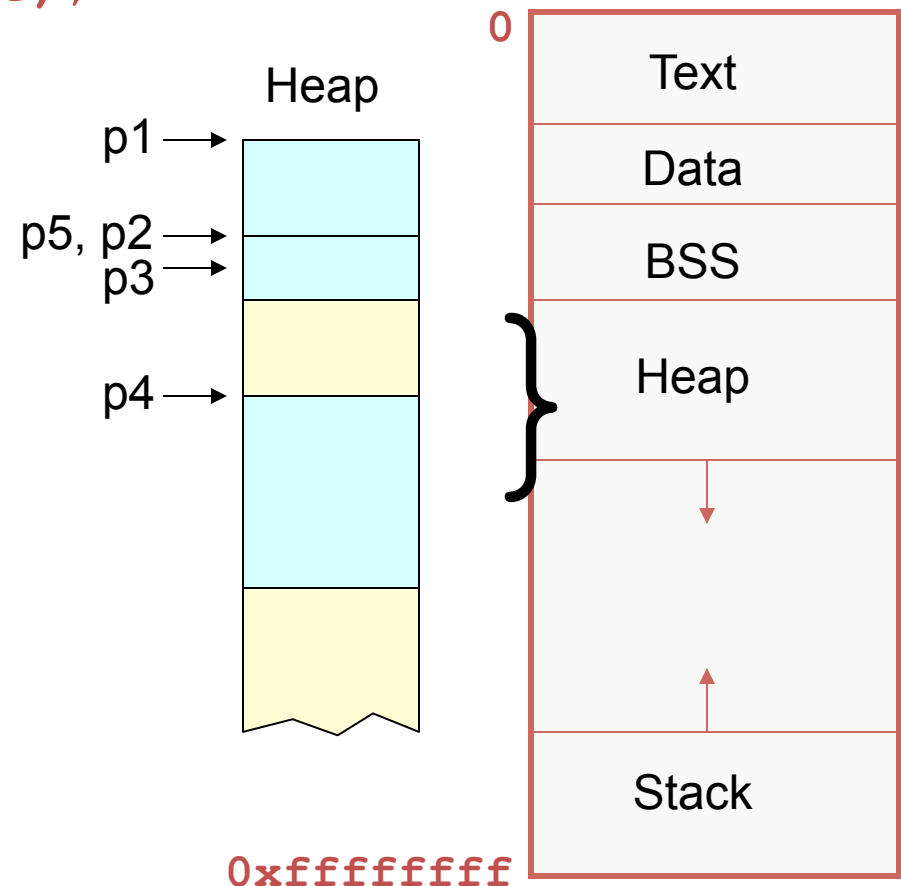
```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
➔ free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Free List (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```

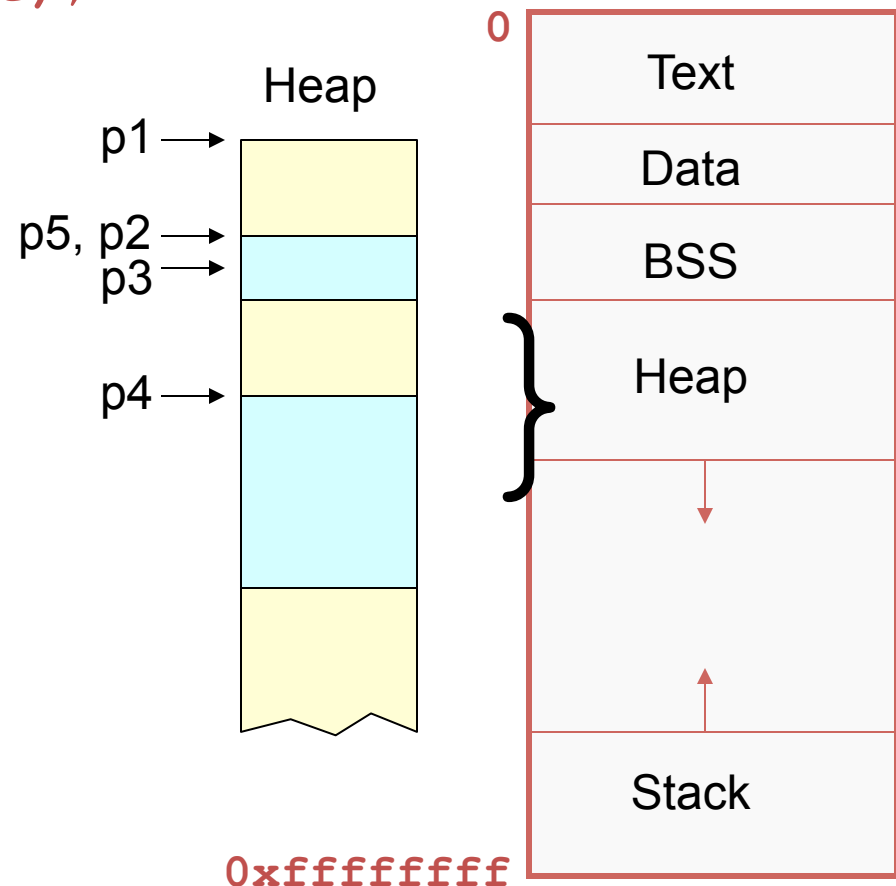
```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
➔ char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Free List (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```


```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
➔ free(p1);
  free(p4);
  free(p5);
```



Free List (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```



```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Free List (C example)

```
#include <stdlib.h>
void *malloc(size_t size);
void free(void *ptr);
```

```
char *p1 = malloc(3);
char *p2 = malloc(1);
char *p3 = malloc(4);
free(p2);
char *p4 = malloc(6);
free(p3);
char *p5 = malloc(2);
free(p1);
free(p4);
free(p5);
```



Free List (C example): Text Description

- Originally free list just holds one big piece of open memory
- List of C commands and their consequences:
 - `char *p1 = malloc(3)`
 - 3 bytes are allocated on the heap
 - p1 holds address 100, the beginning of the 3 bytes
 - `char *p2 = malloc(1)`
 - 1 byte is allocated on the heap
 - p2 holds address 103 the beginning of the byte
 - `char *p3 = malloc(4)`
 - 4 bytes are allocated on the heap
 - p3 holds address 104 the beginning of the 4 bytes
 - `free(p2)`
 - deallocate the memory pointed to by p2
 - now there is a gap of free memory from addresses 103 to 104
 - note: since we are using a free list allocation scheme this memory can be reallocated later
 - free list now has 2 members on it
 - `char *p4 = malloc(6)`
 - look through free list for a piece of memory big enough
 - 6 bytes allocated on the heap
 - p4 holds address 108 the beginning of those 4 bytes
 - `free(p3)`
 - deallocate the memory pointed to by p3
 - now there is a gap of free memory from addresses 103 to 108.
 - the gap from freeing p2 is coalesced with the new gap made by freeing p3
 - `char *p5 = malloc(2)`
 - look through free list for a piece of memory big enough
 - see the gap from 103 to 108 and notice that it can hold 2 bytes
 - split the free memory and allocate p5 in part of it
 - 2 bytes allocated on the heap
 - p4 holds address 103 the beginning of those 2 bytes
 - now is a gap from 105 to 108
 - `free(p1)`
 - deallocate the memory pointed to by p1
 - now is a gap from addresses 100 to 103
 - free list now has 3 members on it
 - `free(p4)`
 - deallocate the memory pointed to by p4
 - gaps in memory are coalesced
 - now free list only has 2 members
 - » a gap from 100 to 103
 - » the rest of the free heap memory
 - `free(p5)`
 - deallocate memory pointed to by p5
 - now all of the original heap memory is free and can be reallocated as needed

What to Do

When You Run Out of Heap


- Ask the operating system for additional memory
 - Ask for a very large chunk of memory
 - ... and insert the new chunk into the free list
 - ... and then try again, this time successfully
- Operating-system dependent
 - E.g., **sbrk** command in UNIX

Performance

- What do we know about the performance of best-fit?
- Slow! Need to scan the free list.
- Trouble: Free chunks are different sizes
- Solution: Binning!
 - Divide list by chunk size

Binning Strategies: Exact Fit

- Have a bin for each chunk size, up to a limit
 - Advantages: no search for requests up to that size
 - Disadvantages: many bins, each storing a pointer
- Except for a final bin for all larger free chunks
 - For allocating larger amounts of memory
 - For splitting to create smaller chunks, when needed




Binning Strategies: Exact Fit

Text Description

- Bins are implemented with linked lists of open memory chunks
- Bin for size 1 has 3 chunks of open memory in it
- Bin for size 2 is empty
- Bin for size 3 has 2 chunks of open memory in it
- Bin for size 4 is empty
- Bin for sizes greater than 4 has 2 chunks of open memory in it
 - one chunk of size 5
 - another chunk of size 8
 - these can be taken and split up later if needed

Binning Strategies: Range

- Have a bin cover a range of sizes, up to a limit
 - Advantages: fewer bins
 - Disadvantages: need to search for a big enough chunk
- Except for a final bin for all larger free chunks
 - For allocating larger amounts of memory
 - For splitting to create smaller chunks, when needed




Binning Strategies: Range

Text Description

- Bin for sizes 1 and 2 holds 3 chunks of open memory
 - 2 chunks of size 1 and 1 chunk of size 2
- Bin for sizes 3 and 4 is empty
- Bin for sizes 4 and 5 holds 2 chunks of open memory
 - 1 chunk of size 4 and the other of size 5
- Bin for sizes 6 and 7 is empty
- Bin for sizes greater than 7 holds 2 chunks of memory
 - One of size 10 and the other of size 14

Deallocation with Free

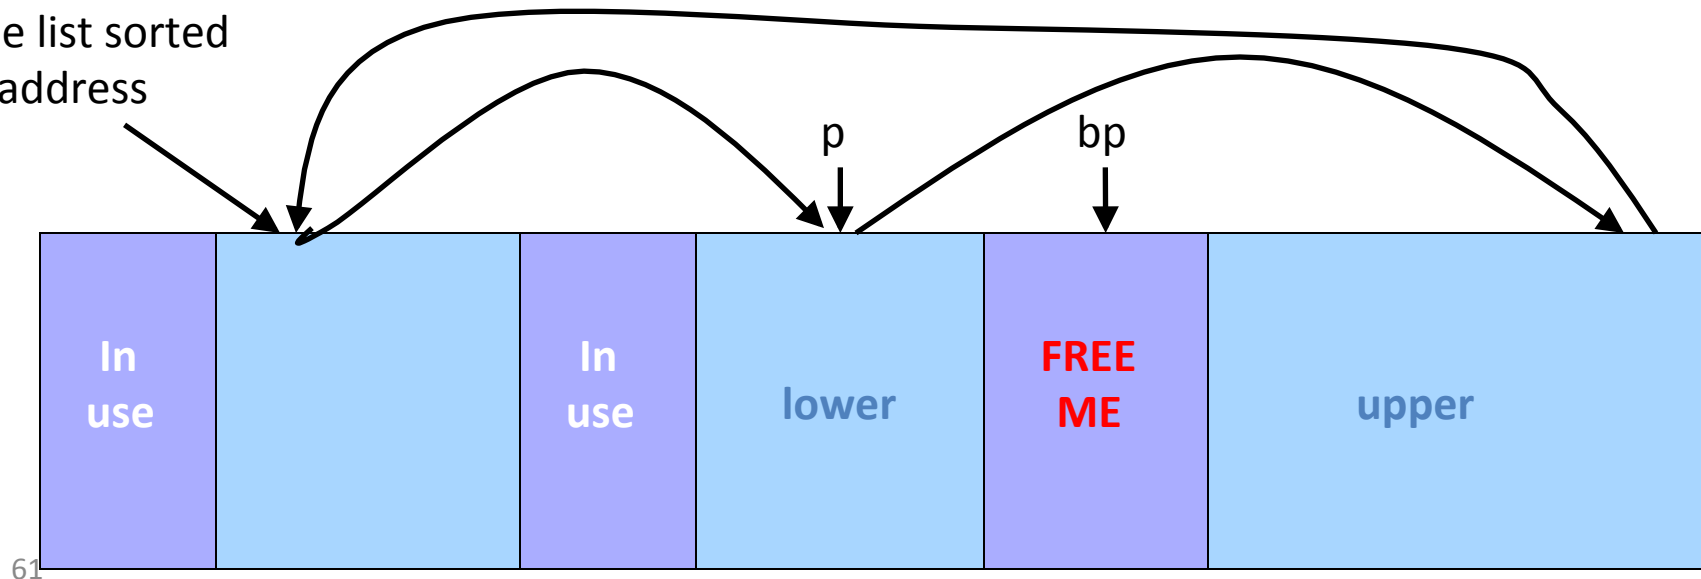
- User passes a pointer to the memory block
- Free function inserts block into the list
 - Identify the start of entry
 - Find the location in the free list
 - Add to the list, coalescing entries, if needed



Coalescing With Neighbors

- Scanning the list finds the location for inserting
 - Pointer to to-be-freed element: **bp**
 - Pointer to previous element in free list: **p**
- Coalescing into larger free blocks
 - Check if contiguous to upper and lower neighbors

Free list sorted
by address



Explicit Memory Management

Challenges for the User

- More code to maintain
- Correctness
 - Free an object too soon -> core dump
 - Free an object too late -> waste space
 - Never free -> at best waste, at worst fail
- Efficiency can be very high
- Gives programmers control

iClicker Question

What advantage does bump pointer allocation have over free-list allocation?

- A. No internal fragmentation
- B. No external fragmentation
- C. Memory re-use
- D. Fast allocation

Summary

- Finished Virtual Memory
- Discussed explicit memory management
 - Allocation policies (bump pointer, free list)
 - De-allocation policies (free)
 - Free-list management

Announcements

- Homework 6 due Friday 8:45a
- Project 2 due Friday, 3/27
- Project 3 posted tonight

Have a good Spring Break! (Be safe.)