

Deadlock Revisited

CS439: Principles of Computer Systems

April 27, 2015

Last Time

Distributed File Systems

- Consistency Models
- NFS
- GFS

Today's Agenda

Deadlocks

- What causes them (again)
- Deadlock Avoidance
- Deadlock Prevention
- Banker's algorithm

Deadlock Revisited

Deadlock, More Formally

- *Deadlock* occurs when two or more threads **or processes** are waiting for an event that can only be generated by these same threads **or processes**
- Deadlock is not starvation
 - Starvation can occur without deadlock
 - occurs when a thread **or process** waits indefinitely for some resources, but other threads **or processes** are actually using it
 - But deadlock does imply starvation

Necessary Conditions for Deadlock

Deadlock *can* happen if all of the following conditions hold:

1. **Bounded Resources:** a finite number of threads **or processes** can use a resource and resources are finite
 - relaxation of mutual exclusion condition
2. **Hold and Wait:** at least one thread **or process** holds a resources and is waiting for other resources to become available. A different thread holds the resource.
3. **No Pre-emption:** a thread **or process** only releases a resource voluntarily; another thread, **process**, or the OS cannot force the thread **or process** to release the resource
4. **Circular Wait:** A set of waiting **processes or threads** $\{t_1, \dots, t_n\}$ where t_i is waiting on t_{i+1} ($i=1$ to n) and t_n is waiting on t_1

Managing Deadlocks

- **Deadlock prevention** adopts a policy that breaks one of the four conditions
- **Deadlock avoidance** algorithms check resource requests and possible availability to prevent deadlock
 - Guarantee that deadlock will never occur
 - Breaks one of the four necessary conditions
- **Deadlock detection** algorithms find instances of deadlock and try to recover
 - Admit the possibility of deadlock occurring and periodically check for it

Deadlock Prevention

Prevent deadlock by insuring that at least one of the necessary conditions doesn't hold

1. **Bounded Resources:** make resources sharable or provide more resources
2. **Hold and Wait:** guarantee a thread **or process** cannot hold one resource when it requests another (or must request all at once)
3. **No Pre-emption:** If a thread **or process** requests a resource that cannot be immediately allocated to it, then the OS pre-empts all the resources the thread **or process** is currently holding. Only when all the resources are available will the OS restart the thread **or process**
4. **Circular Wait:** Impose an ordering on the resources and request them in order

Deadlock Prevention: Resource Ordering

- Order all locks (or semaphores or resources)
- All code grabs locks in a predefined order
- Complications:
 - Maintaining global order is difficult in a large project
 - Global order can force a client to grab a lock earlier than it would like, tying up a resource for longer than necessary
- What happens when we apply this to system resources?

Avoiding Deadlock: The Banker's Algorithm



- Allows sum of maximum resource needs to exceed the total available resources
as long as there exists a schedule of loan fulfillments such that all clients can:
 - Receive their maximal loan
 - Build their respective houses
 - Pay back all the loan
- More efficient than atomically acquiring all resources

Avoiding Deadlock:

The Banker's Algorithm


Plain Text

- Allows sum of maximum resource needs to exceed the total available resources as long as there exists a schedule of loan fulfillments such that:
 - All clients receive their maximal loan
 - Build their respective houses
 - pay back all the loan
- More efficient than atomically acquiring all resources
- Picture is from the movie *It's a Wonderful Life*, a classic starring Jimmy Stewart. In that movie, he is a banker.

The Banker's Algorithm: Details

- Banker has N units, but loans out many more
 - Okay as long as $N+1$ units are not needed at the same time
- Uses *safe* and *unsafe* states
 - *Safe* states are states where enough resources are potentially available such that at least one process can run to completion
 - *Unsafe* states may lead to deadlock
- If resource request leads to an unsafe state, request is denied even if resources are currently available

Living Dangerously: Safe, Unsafe, Deadlocked



A system's trajectory
through its state space

- **Safe:** For any possible set of resource requests, there exists one safe schedule of processing requests that succeeds in granting all pending and future requests
 - no deadlock as long as system can enforce safe schedule
- **Unsafe:** There exists a set of (pending and future) resource requests that leads to a deadlock, for any schedule in which requests are processed
 - unlucky set of requests can force deadlock
- **Deadlocked:** The system has at least one deadlock

Living Dangerously: Safe, Unsafe, and Deadlocked

Plain Text

- Safe: for any possible set of resource requests, there exists one safe schedule of processing requests that succeeds in granting all pending and future requests
 - No deadlock as long as system can enforce safe schedule
- Unsafe: there exists a set of (pending and future) resource requests that leads to a deadlock, for any schedule in which requests are processed
 - Unlucky set of requests can force deadlock
- Deadlocked: the system has at least one deadlock
- A system's trajectory through its state space (as the states relate to deadlock)
 - Safe and deadlocked states are completely disjoint
 - Must go from unsafe to deadlocked state (cannot go straight to deadlock from safe)
 - Can move between safe and unsafe states without ever going into deadlock

Banker's Algorithm: Example

- 5 processes, 4 resources

Max				
	R ₁	R ₂	R ₃	R ₄
P ₁	0	0	1	2
P ₂	1	7	5	0
P ₃	2	3	5	6
P ₄	0	6	5	2
P ₅	0	6	5	6

Allocated				
	R ₁	R ₂	R ₃	R ₄
P ₁	0	0	1	2
P ₂	1	0	0	0
P ₃	1	3	5	3
P ₄	0	6	3	2
P ₅	0	0	1	4

Available (to be allocated)			
R ₁	R ₂	R ₃	R ₄
1	5	2	0

- Is this a safe state?

Banker's Algorithm: Example

Plain Text

- 5 processes (represented by P1-P5) competing for 4 resources (represented by R1-R4)
- First metric: maximum number of each kind of resource each process will need
 - P1: 0 R1s, 0 R2s, 1 R3, and 2 R4s
 - P2: 1 R1, 7 R2s, 5 R3s, and 0 R4s
 - P3: 2 R1s, 3 R2s, 5 R3s, and 6 R4s
 - P4: 0 R1s, 6 R2s, 5 R3s and 2 R4s
 - P5: 0 R1s, 6 R2s, 5 R3s and 6 R4s
- Second metric: number of each kind of resource that has already been allocated to each process
 - P1: 0 R1s, 0 R2s, 1 R3, and 2 R4s
 - P2: 1 R1, 0 R2s, 0 R3s, and 0 R4s
 - P3: 1 R1, 3 R2s, 5 R3s, and 3 R4s
 - P4: 0 R1s, 6 R2s, 3 R3s, and 2 R4s
 - P5: 0 R1s, 0 R2s, 1 R3, and 4 R4s
- Third metric: number of each kind of resource in the system that is still available to be allocated
 - R1: 1
 - R2: 5
 - R3: 2
 - R4: 0
- Big question: is this a safe state?

Example: Determining Safety

- 5 processes, 4 resources

Max					Allocated					Available					Max Request				
	R ₁	R ₂	R ₃	R ₄		R ₁	R ₂	R ₃	R ₄		R ₁	R ₂	R ₃	R ₄		R ₁	R ₂	R ₃	R ₄
P ₁	0	0	1	2	-	P ₁	0	0	1	2	=	P ₁	0	0	0	0			
P ₂	1	7	5	0		P ₂	1	0	0	0		P ₂	0	7	5	0			
P ₃	2	3	5	6		P ₃	1	3	5	3		P ₃	1	0	0	3			
P ₄	0	6	5	2		P ₄	0	6	3	2		P ₄	0	0	2	0			
P ₅	0	6	5	6		P ₅	0	0	1	4		P ₅	0	6	4	2			

- Determine MaxRequest by subtracting Allocated from Maximum

Example: Determining Safety

- 5 processes, 4 resources

Max					Allocated					Available				Max Request						
	R ₁	R ₂	R ₃	R ₄		R ₁	R ₂	R ₃	R ₄		R ₁	R ₂	R ₃	R ₄		R ₁	R ₂	R ₃	R ₄	
P ₁	0	0	1	2	-	P ₁	0	0	1	2	=	1	5	2	0	P ₁	0	0	0	0
P ₂	1	7	5	0		P ₂	1	0	0	0		P ₂	0	7	5	0				
P ₃	2	3	5	6		P ₃	1	3	5	3		P ₃	1	0	0	3				
P ₄	0	6	5	2		P ₄	0	6	3	2		P ₄	0	0	2	0				
P ₅	0	6	5	6		P ₅	0	0	1	4		P ₅	0	6	4	2				

- While safe sequence does not include all processes:
 - Is there a P_i such that $\text{MaxRequest}_i \leq \text{Available}$?
 - if no, exit with unsafe
 - if yes, add P_i to the sequence and set $\text{Available} = \text{Available} + \text{Allocated}$

Example: Determining Safety

- 5 processes, 4 resources

Max				
	R ₁	R ₂	R ₃	R ₄
P ₁	0	0	1	2
P ₂	1	7	5	0
P ₃	2	3	5	6
P ₄	0	6	5	2
P ₅	0	6	5	6

Allocated				
	R ₁	R ₂	R ₃	R ₄
P ₁	0	0	1	2
P ₂	1	0	0	0
P ₃	1	3	5	3
P ₄	0	6	3	2
P ₅	0	0	1	4

Available			
R ₁	R ₂	R ₃	R ₄
1	5	2	0

- Is this a safe state? YES!

Example: Determining Safety

Plain Text

- Using the same metrics from slide 15
 - 5 processes competing for 4 resources
- Determine maximum number of resources that any process may still request by subtracting allocated from maximum
- Max requests by process:
 - P1:
 - Max needed: 0 R1s, 0 R2s, 1 R3, and 2 R4s
 - Allocated: 0 R1s, 0 R2s, 1 R3, and 2 R4s
 - Max request = max needed - allocated = 0 R1s, 0 R2s, 0 R3s, and 0 R4s
 - P2:
 - Max needed: 1 R1, 7 R2s, 5 R3s, and 0 R4s
 - Allocated: 1 R1, 0 R2s, 0 R3s, and 0 R4s
 - Max request = max needed - allocated = 0 R1s, 7 R2s, 5 R3s, and 0 R4s
 - P3:
 - Max needed: 2 R1s, 3 R2s, 5 R3s, and 6 R4s
 - Allocated: 1 R1, 3 R2s, 5 R3s, and 3 R4s
 - Max request = max needed - allocated = 1 R1, 0 R2s, 0 R3s, and 3 R4s
 - P4:
 - Max needed: 0 R1s, 6 R2s, 5 R3s, and 2 R4s
 - Allocated: 0 R1s, 6 R2s, 3 R3s, and 2 R4s
 - Max request = max needed - allocated = 0 R1s, 0 R2s, 2 R3s, and 0 R4s
 - P5:
 - Max needed: 0 R1s, 6 R2s, 5 R3s, and 6 R4s
 - Allocated: 0 R1s, 0 R2s, 1 R3, and 4 R4s
 - Max request = max needed - allocated = 0 R1s, 6 R2s, 4 R3s, and 2 R4s
- While safe sequence does not include all processes:
 - Is there P_i such that $\text{MaxRequest}_i \leq \text{available}$?
 - if no, exit with unsafe
 - Test this by creating a safe sequence that allows each process to complete.
- Is the state described before safe? YES!
 - We are able to create a sequence that allows each process to finish.

Updated Example: Determining Safety

- 5 processes, 4 resources

Max				
	R ₁	R ₂	R ₃	R ₄
P ₁	0	0	1	2
P ₂	1	7	5	0
P ₃	2	3	5	6
P ₄	0	6	5	2
P ₅	0	6	5	6

Allocated				
	R ₁	R ₂	R ₃	R ₄
P ₁	0	0	1	2
P ₂	1	0	0	0
P ₃	1	3	5	3
P ₄	0	6	3	2
P ₅	0	0	1	4

Available			
R ₁	R ₂	R ₃	R ₄
1	5	2	0

- P2 wants to change its allocation to

0	4	2	0
---	---	---	---
- Safe?

Updated Example: Determining Safety

- 5 processes, 4 resources

Max					Allocated					Available				Max Request				
	R ₁	R ₂	R ₃	R ₄		R ₁	R ₂	R ₃	R ₄	R ₁	R ₂	R ₃	R ₄		R ₁	R ₂	R ₃	R ₄
P ₁	0	0	1	2		0	0	1	2	2	1	0	0	P ₁	0	0	0	0
P ₂	1	7	5	0		0	4	2	0					P ₂	1	3	3	0
P ₃	2	3	5	6		1	3	5	3					P ₃	1	0	0	3
P ₄	0	6	5	2		0	6	3	2					P ₄	0	0	2	0
P ₅	0	6	5	6		0	0	1	4					P ₅	0	6	4	2

- P2 wants to change its allocation to

0	4	2	0
---	---	---	---
- Safe? No!

Updated Example: Determining Safety

Plain Text

- P2 wants to change its allocated to: 0 R1s, 4 R2s, 2 R3s and 0 R4s
- This would change the metrics to:
 - First metric: maximum number of each kind of resource each process will need
 - P1: 0 R1s, 0 R2s, 1 R3, and 2 R4s
 - P2: 1 R1, 7 R2s, 5 R3s, and 0 R4s
 - P3: 2 R1s, 3 R2s, 5 R3s, and 6 R4s
 - P4: 0 R1s, 6 R2s, 5 R3s, and 2 R4s
 - P5: 0 R1s, 6 R2s, 5 R3s, and 6 R4s
 - Second metric: number of each kind of resource that has already been allocated to each process
 - P1: 0 R1s, 0 R2s, 1 R3, and 2 R4s
 - P2: 0 R1s, 4 R2s, 2 R3s, and 0 R4s
 - P3: 1 R1, 3 R2s, 5 R3s, and 3 R4s
 - P4: 0 R1s, 6 R2s, 3 R3s, and 2 R4s
 - P5: 0 R1s, 0 R2s, 1 R3, and 4 R4s
 - Third metric: number of each kind of resource in the system that is still available to be allocated
 - R1: 2
 - R2: 1
 - R3: 0
 - R4: 0
- Updated max requests by process:
 - P1:
 - Max needed: 0 R1s, 0 R2s, 1 R3, and 2 R4s
 - Allocated: 0 R1s, 0 R2s, 1 R3, and 2 R4s
 - Max request = max needed - allocated = 0 R1s, 0 R2s, 0 R3, and 0 R4s
 - P2:
 - Max needed: 1 R1, 7 R2s, 5 R3s, and 0 R4s
 - Allocated: 0 R1, 4 R2s, 2 R3s, and 0 R4s
 - Max request = max needed - allocated = 1 R1, 3 R2s, 3 R3s, and 0 R4s
 - P3:
 - Max needed: 2 R1s, 3 R2s, 5 R3s, and 6 R4s
 - Allocated: 1 R1, 3 R2s, 5 R3s, and 3 R4s
 - Max request = max needed - allocated = 1 R1, 0 R2s, 0 R3s, and 3 R4s
 - P4:
 - Max needed: 0 R1s, 6 R2s, 5 R3s, and 2 R4s
 - Allocated: 0 R1s, 6 R2s, 3 R3s, and 2 R4s
 - Max request = max needed - allocated = 0 R1s, 0 R2s, 2 R3s, and 0 R4s
 - P5:
 - Max needed: 0 R1s, 6 R2s, 5 R3s, and 6 R4s
 - Allocated: 0 R1s, 0 R2s, 1 R3, and 4 R4s
 - Max request = max needed - allocated = 0 R1s, 6 R2s, 4 R3s, and 2 R4s
- Is this allocation safe? NO!
 - We cannot create a sequence of requests such that all processes can finish.

Detecting Deadlock: Work at Home Problem

- 5 processes, 3 resources

Allocated			
	R ₁	R ₂	R ₃
P ₁	0	1	0
P ₂	2	0	0
P ₃	3	0	3
P ₄	2	1	1
P ₅	0	0	2

Available		
R ₁	R ₂	R ₃
0	0	0

Max Request			
	R ₁	R ₂	R ₃
P ₁	0	0	0
P ₂	2	0	2
P ₃	0	0	0
P ₄	1	0	2
P ₅	0	0	2

- Given the set of pending requests is there a safe sequence?
 - If no, then deadlock!

Detecting Deadlock: Work at Home Problem

- 5 processes, 3 resources

Allocated			
	R ₁	R ₂	R ₃
P ₁	0	1	0
P ₂	2	0	0
P ₃	3	0	3
P ₄	2	1	1
P ₅	0	0	2

Available		
R ₁	R ₂	R ₃
0	0	0

Max Request			
	R ₁	R ₂	R ₃
P ₁	0	0	0
P ₂	2	0	2
P ₃	0	0	1
P ₄	1	0	2
P ₅	0	0	2

- Given the set of maximum requests is there a safe sequence?
 - If no, then deadlock!

Detecting Deadlock: Work at Home Problem Plain Text

- 5 processes competing for 3 resources
- Allocated resources:
 - P1: 0 R1s, 1 R2, and 0 R3s
 - P2: 2 R1s, 0 R2s, and 0 R3s
 - P3: 3 R1s, 0 R2s, and 3 R3s
 - P4: 2 R1s, 1 R2, and 1 R3
 - P5: 0 R1s, 0 R2s, and 2 R3s
- Available resources:
 - R1: 0
 - R2: 0
 - R3: 0
- Max requests:
 - P1: 0 R1s, 0 R2s, and 0 R3s
 - P2: 2 R1s, 0 R2s, and 2 R3s
 - P3: 0 R1s, 0 R2s, and 0 R3s
 - P4: 1 R1, 0 R2s, and 2 R3
 - P5: 0 R1s, 0 R2s, and 2 R3s
- Given the set of maximum requests is there a safe sequence?
 - If no, the deadlock!
- Now consider a change in the system such that P3's max request becomes P3: 0 R1s, 0 R2s, and 1 R3
 - Is there a safe sequence?

iClicker Question

The Banker's Algorithm is a good choice for deadlock avoidance in a modern day OS

- A. True
- B. False

Weaknesses of Deadlock Avoidance Algorithms

- Must know resource requests of processes up front
- Processes may not enter the system
- Resources are assumed to always be working

Deadlock Detection

Resource Allocation Graphs

- Loosely, graphs the state of resources in the system
- Used to detect deadlock
- Threads/processes are represented by circles
- Resources are represented by squares
- Arrows represent dependency
 - Arrows from a thread to a resource indicate “waiting for”
 - Arrows from resources to threads indicate “owned by”


Resource Allocation Graphs, Formally

- Basic components of any resource allocation problem
 - Processes and resources
- Model the state of a computer system as a directed graph
 - $G = (V, E)$
 - $V =$ the set of vertices $= \{P_1, \dots, P_n\} \cup \{R_1, \dots, R_m\}$



$E =$ the set of edges $=$

$\{\text{edges from a resource to a process}\} \cup \{\text{edges from a process to a resource}\}$



Resource Allocation Graphs, Formally:


Plain Text

- Basic components of any resource allocation problem
 - Processes and resources
- Model the state of a computer system as a directed graph
 - $G = (V, E)$
 - $V = \text{set of vertices} = \{P_1, \dots, P_n\} \text{ union } \{R_1, \dots, R_m\}$
 - $E = \text{the set of edges} = \{\text{edges from a resource to a process}\} \text{ union } \{\text{edges from a process to a resource}\}$
 - An edge from a process to a resource exists if the process is requesting that resource
 - An edge from a resource to a process exists if that resource is allocated to that process
- Resources modeled as rectangles that contain marks for every copy of said resources
 - ex: If there were multiple copiers there would just be 1 copier vertex with a mark for each copier that exists

Resource Allocation Graphs: An Example

A PostScript interpreter that is waiting for the frame buffer lock and a visualization process that is waiting for the memory frames lock

$$V = \{PS\ interpret, visualization\} \cup \{memory\ frames, frame\ buffer\ lock\}$$



Resource Allocation Graphs:


An Example

Plain Text

- $V = \{\text{Processes: PS interpreter, visualization}\} \cup \{\text{Resources: memory frames, frame buffer lock}\}$
- Frame buffer has been allocated to the visualization process
- All memory frames have been allocated to the PostScript interpreter
- A PostScript interpreter that is waiting for the frame buffer lock and a visualization process that is waiting for memory
- The graph:
 - Nodes: memory frames (R1), PostScript Interpreter (P), Frame Buffer (R2) and Visualization Process (P)
 - Edges:
 - All 3 memory frames to PostScript Interpreter
 - PostScript Interpreter to frame buffer
 - Frame buffer to Visualization Process
 - Visualization Process to Memory Frames
 - There are cycles in the graph! What does this mean?

Resource Allocation Graphs: Cycles


Theorem: If a resource allocation graph does not contain a cycle, then no processes are deadlocked.



Resource Allocation Graphs: Cycles

A cycle in a *RAG* is a necessary condition for deadlock

Is the existence of a cycle a sufficient condition?



Resource Allocation Graphs: Cycles

Plain Text

- Theorem: if a resource allocation graph does not contain a cycle, then no processes are deadlocked
- A cycle in a RAG is necessary condition for deadlock
- Is the existence of a cycle a sufficient condition?
 - NO!
- Example in update of RAG from previous slide
 - Nodes: memory frames (R1), PostScript Interpreter (P), Frame Buffer (R2) and Visualization Process (P), Game (P)
 - Edges:
 - 2 memory frames allocated to PostScript Interpreter (means there is still 1 available)
 - 1 memory frame allocated to the game
 - PostScript Interpreter requesting frame buffer
 - Frame buffer allocated to Visualization Process
 - Visualization Process requesting memory frames
 - Over the course of time, the game process releases its memory frame
- How does this update change things?
 - Now the visualization process could get a memory frame and make progress

Cycles, In Words

- If the graph has no cycles, no deadlock exists
- If the graph has a cycle, deadlock *might* exist
 - If there is only a single unit of all resources then a set of processes are deadlocked if and only if there is a cycle in the resource allocation graph
 - If there are multiple instances of a resource(s) and any instance of a resource involved in the cycle is held by a thread/process that is not in the cycle, progress might be made when that thread/process releases the resource

Deadlock Detection

Using a resource allocation graph, scan the graph for cycles, then break the cycles

- Kill threads or processes in the cycle
- Kill threads or processes one at a time, forcing each to give up its resources
- Pre-empt resources one at a time, and rollback the state of the thread/process holding the resource to its state prior to acquiring the resource
 - Common to database transactions

Deadlock Detection

- Detecting cycles in the graph requires $O(n^2)$, where n is the number of vertices (processes and resources)
- When should we execute the algorithm?
 - Just before granting a resource?
 - When a request is denied?
 - On a regular schedule?
 - When CPU utilization drops below some threshold?

Deadlock Handling: Real Life

- Ostrich Algorithm

Summary

Deadlock is a situation in which a set of threads/processes cannot proceed because each requires resources held by another member of the set. Approaches to handling deadlock are:

- Prevention: design resource allocation strategies that guarantee that one of the necessary conditions never holds
- Avoidance: don't allocate a resource if it would introduce a cycle
- Detection and recovery: recognize deadlock after it has occurred and break it
- In real life:
 - For resources managed by the program, code carefully! (Does not work for OS managed resources)
 - Ignore the possibility!

Announcements

- Homework 11 (Last one!) due Friday, 5/1, 8:45a
- Project 4 (Last one!) due Friday, 5/8, 11:59p
 - No slip days!
- If you have a conflict for the final, you should have already contacted me (email, please!)
 - Thursday, May 14, 7p-10p in UTC 2.102A