

# CS439: Principles of Computer Systems

Dr. Alison N. Norman  
Department of Computer Science  
The University of Texas at Austin  
Spring 2015

# Who am I?

- Education
  - Undergrad in CS from Georgia Tech
  - MS and Ph.D. in CS from UT Austin
- Research
  - Supercomputing
- Family
  - Married
  - Three children: two boys (7 and 3) and one girl (8m)
  - Two dogs

# Today's Plan

- Introduce and motivate course themes
- Course organization and logistics
- Quiz

# Why are we here?

Two main goals:

- Learn the low-level software abstractions that make the computer work
  - Operating System
  - Network
  - Various aspects of memory management
- Use these topics as a case study to understand large-scale system design

# System Design

How do we construct systems that are

- reliable
- portable
- efficient
- secure

?

# What is an OS?

- No universally accepted definition
  - Is it everything that comes on a computer?
    - Used to be, then came Microsoft (US v. Microsoft, 1998)
    - Now this varies widely
- Program that is always running
  - Ha.

# Operating System: A definition


Software that manages a computer's resources

- makes it easier to write the applications you want to write
- makes you want to use the applications you wrote by running them efficiently

# Why Study Operating Systems?

- To learn how computers work
- To learn how to manage complexity through appropriate abstractions
- To learn about system design
  - Performance vs. simplicity, HW vs SW, etc
  - Design trade-offs made in the past do not necessarily apply now
  - Those made now will not necessarily apply in the future
- Operating Systems are everywhere!





Where's the Operating System?  
Las Vegas!



Where's the Operating System?  
New York!

# Operating Systems: More than One Hat

- Referee
  - Manages shared resources
- Illusionist
  - Infinite memory! Your own private processor!
- Glue
  - Provides standard services which the hardware implements

# Operating Systems as Referee

- Resource allocation
  - Coordinates multiple applications and users to achieve fairness and efficiency
- Isolation
  - Protects processes from one another
    - One application's bugs should not crash another (or the whole system!)
    - If it does crash, should fail gracefully
- Communication
  - Allow processes to work together

# Operating Systems as Illusionist

Illusion of resources that are not really present

- Virtualization: processor, memory, screen space
- Entire computer!

# Operating Systems as Glue

Provides standard services to simplify application design and facilitate sharing

- File system, virtual memory, networking
- Decouples hardware and application development
- Start, stop, and clean up after a program

# Evaluating an Operating System

- Reliability
  - OS does exactly what is designed to do
- Security
  - OS cannot be compromised by a malicious attacker
- Portability
  - OS does not change as hardware changes
- Performance
  - efficiency, overhead, fairness, latency, throughput, predictability

# Reliability

- The ability of a computer-related hardware or software component to consistently perform according to its specifications.
- In theory, a reliable product is totally free of technical errors (yeah, right)
- Availability: percentage of time system is useful
  - Depends on *MTTF* and *MTTR*




# Security

- Includes privacy: data on the computer only accessible to authorized users
- Strong fault isolation helps, but not enough
  - Security mechanisms should not prevent legitimate sharing!
- Enforcement mechanism
  - Ensures only permitted actions are allowed
- Security policy
  - Defines what is permitted

# Portability

- OSs can live longer than your cat!
  - must support applications not yet written
  - must run on hardware not yet developed
- Three interfaces
  - Abstract Machine Interface (AMI)
    - between OS and apps: API + memory access model + legally executable instructions
  - Application Programming Interface (API)
    - function calls provided to apps
  - Hardware Abstraction Layer (HAL)
    - abstracts hardware *internally to the OS*

# Logical OS Structure



# Performance

- Efficiency/Overhead
  - how much is lost by not running on bare hardware?
- Fairness
  - how are resources divided?
- Response time
  - how long does a task take to complete
- Throughput
  - how many tasks complete per unit of time
- Predictability
  - are performance metrics consistent over time?

# What You'll Learn in this Course

## 1. How to approach problems

- Fundamental issues
- Design space
- Manage complexity
- Case studies

Goal: You will be able to devise good solutions to similar (and very different) problems.

# What You'll Learn in this Course

2. Specific techniques you should be able to apply to other problems
  - Time-tested solutions to hard problems
  - Goal: be a good engineer
3. Details(ish) of modern operating systems
  - Lots of material, changes quickly
  - Not a priority of this class
    - Would rather you know the abstractions so that you can apply your knowledge to the next OS

# Things You'll Encounter

- Design Problems
  - Understand the problem and define it
  - Understand the space of possible solutions and previous approaches
  - Formulate your own approach and justify it
- Implementation Issues
  - Real systems are more difficult to build than explain
  - The devil is in the details

# Summary

- Operating Systems are infinite loops that manage resources
- Key ideas: coordination and abstraction
- It's Going to Be Great!



# Course Organization and Logistics

# What knowledge you need to begin this course

- Prerequisites: CS429(H) with a grade of at least a C-
- Solid basic understanding of hardware
- Solid programming skills (especially in C)

You must understand the components to understand the implications of how they interact!

# Teaching Staff

- Teaching Assistants:
  - Soumyajit Gupta
  - Alex Knaust
  - Megan Avery
  - Jeremy Hintz
  - James Lee
  - Jon Lee
  - Robert Lynch
  - Sage Moore
  - Ben Ridings
- And, obviously, me.

# Course Materials

- Website: Go-to place for information
  - Syllabus, Schedule, Projects, Homeworks, Slides, Useful links, Feedback form
- Textbooks:
  - *Computer Systems: A Programmer's Perspective* by Bryant & O'Hallaron (from 429)
  - *Operating Systems: Three Easy Pieces* by Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau (online, follow links on syllabus or schedule)
  - *Operating Systems and Middleware: Supporting Controlled Interaction* by Max Hailperin (online, follow links on syllabus or schedule)
- Piazza: discussion board
  - Course "CS439N"
  - Many of you received an invitation
- Canvas: grade center
  - Please ignore Canvas's calculation of your grades
- iClicker: participation counts, so get one and get it registered!
  - Participation points begin next week

# Schedule Overview

- Introduction
- Concurrency and Synchronization
- Memory Management
- File Systems: Use and Implementation
- Networked Systems
- Parallel and Distributed Computing (briefly)
- Security (briefly)

# Each Class

- Introduces concepts, covers high-level ideas
- Timing:
  - Approximately 50 minutes of lecture
  - Approximately 5 minutes of break
  - Approximately 45 minutes of lecture

# Discussion Sections

- Required!
- You MUST attend your own
- First half will be about lecture and homework
  - an additional homework question
  - then discussion of that week's homework
- Second half will be discussion of current project

# Homeworks

- Weekly homeworks (eleven total)
- Designed to help you prepare for the exams
- All but one question will be posted online and due at 8:45am on Fridays
- One question will be solved and turned in during discussion section
- **Must turn in both sets of problems to receive credit for that homework**
- Graded on an okay/not okay basis (Binary!)
- Lowest two dropped
  - These are your excused absences



# Projects

- There will be 5 projects in this course
- Pair or group programming
- Each project is accompanied by a design document
- They will not be equally weighted
- They will be difficult
  - Systems programming is difficult!
- Your life will be easier if you learn the Linux environment

# Expected Effort

- This is a hard course that requires a LOT of effort
- Topics are new and detailed
- There are many design tradeoffs to understand
- Systems programming is hard
  - Debugging systems code is worse
- Projects can take 10-15 hours in the beginning, and 30-40 hours later in the semester
  - *If* it goes well
  - Start early, stay late

# Evaluation

- Projects (32%)
  - Build operating system components
  - 4 slip days total
    - 2 maximum on each project
    - None on last project
  - Due 11:59pm on select Fridays
  - More information soon
- Homeworks (8%)
  - Written
  - Due 8:45a Fridays *and* at the beginning of discussion section
  - Graded on an *ok/not okay* basis

# Evaluation

- Exams and a final (16%, 16%, 22%)
  - Exams are 2/25 and 4/8 (mark your calendars!)
    - In the evening. Locations are on the schedule.
  - Final is as scheduled by the registrar
    - Will NOT be at time currently listed
- iClicker participation (6%)
  - Instant feedback for me
  - Gives you a reason to come to class
  - Attend 80% of the classes for full credit
    - The other 20% is for “excused absences”.
  - Using laptops or other digital devices forfeits your participation credit
- Final grades will be curved (a little)
  - If you are on the edge, you need to have shown effort
    - Attended class
    - Turned in all assignments
  - Or I will NOT bump you up
- Extra credit: 10 points on final exam for reading 75% of Piazza posts (as measured by Piazza)

# Collaboration and Cheating

- Collaboration
  - Discuss problem sets and programming assignments
  - Discuss possible interpretation of questions, technical details
- Cheating
  - Copying solutions code or programs from someone else, previous semesters' solutions, or public domain
  - Providing material for someone else to imitate
  - Participating in discussion group where one person writes solution and everyone else copies it
  - **Penalty for cheating is an F in the course and a referral to the Dean of Students office**

# How to Succeed in This Course

- Keep up
- Attend class
- Do the reading
- Do the projects (and start them early!)
- Ask questions
- Get to know the people in the class
  - How many people you know is the number one indicator of success
    - study groups, problem discussion, etc.

# How to Get Help

- Ask questions!
  - In class
  - Office Hours
    - In online syllabus and on Office Hours webpage
    - *All begin Monday, 1/26*
  - Piazza
    - Use Anonymous feature if necessary
- Many resources/tutorials on course main page
  - Online Lectures, C tutorials, link to B&O student site,...

# Other Thoughts

- Enrollment is high
- Workload is heavy
- Grading will be slow
- Use of discussion board essential



# C and Linux

# C and Linux

- This course relies heavily on C and Linux
- You should have prior knowledge of these from 429
- Your first discussion section (next Friday!) will provide introduction and review (among other topics)
- I want to know where you are now so we can plan that review

# Learning C and Linux

- Resources on the website
  - I'll add more
- Next Thursday's discussion section will review C and Linux
- ASK Questions
  - Early and often

# Assessment

When you are finished, bring your paper up to me and you may leave.

See you Monday!