

CS 33

Machine Programming (3)

Many of the slides in this lecture are either from or adapted from slides provided by the authors of the textbook “Computer Systems: A Programmer’s Perspective,” 2nd Edition and are provided from the website of Carnegie-Mellon University, course 15-213, taught by Randy Bryant and David O’Hallaron in Fall 2010. These slides are indicated “Supplied by CMU” in the notes section of the slides.

Stacks and Functions (1)

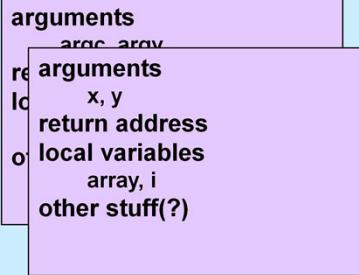
```
int main(int argc, char *argv[]) {  
    char buf[1024];  
    int res, a1, a2, a3, a4;  
    ...  
    res = sub1(a1, a2);  
    a3 = res +1;  
    ...  
    res = sub2(a2, a3, a4);  
    res += 6;  
    ...  
    return 0;  
}
```

arguments
 argc, argv
return address
local variables
 buf, res, a1, a2, a3, a4
other stuff(?)

main's "context"

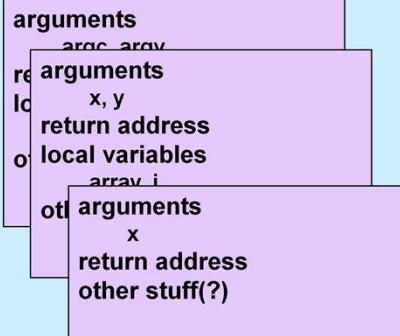
Stacks and Functions (2)

```
int sub1(int x, int y) {  
    int array[24];  
    int i;  
  
    ...  
  
    for (i=0; i<24; i++)  
        array[i] = fib[i];  
  
    ...  
  
    return sub10(24, array);  
}
```



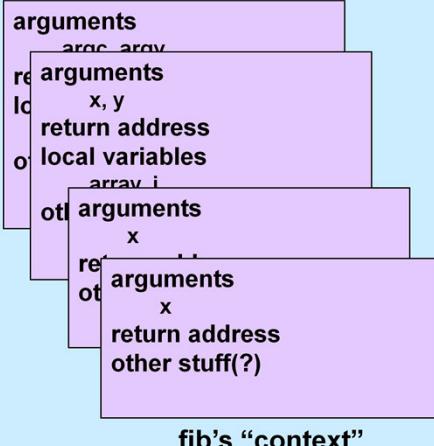
Stacks and Functions (3)

```
int fib(int x) {  
    if (x < 2)  
        return x;  
    else  
        return fib(x-1)+fib(x-2);  
}
```



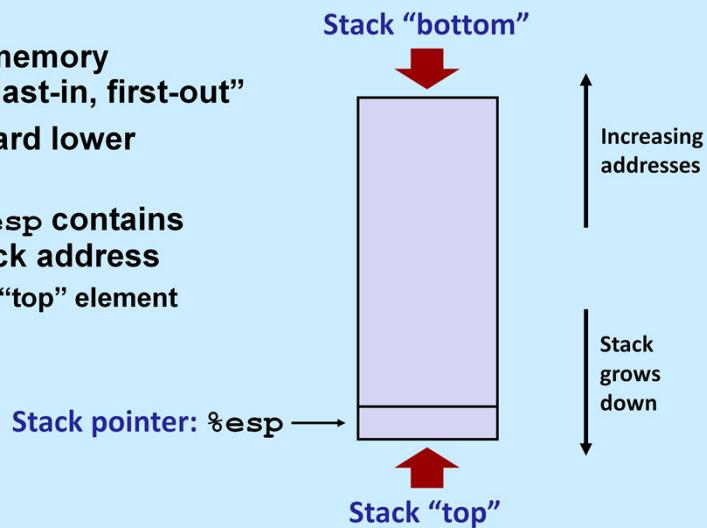
Stacks and Functions (4)

```
int fib(int x) {  
    if (x < 2)  
        return x;  
    else  
        return fib(x-1)+fib(x-2);  
}
```



IA32 Stack

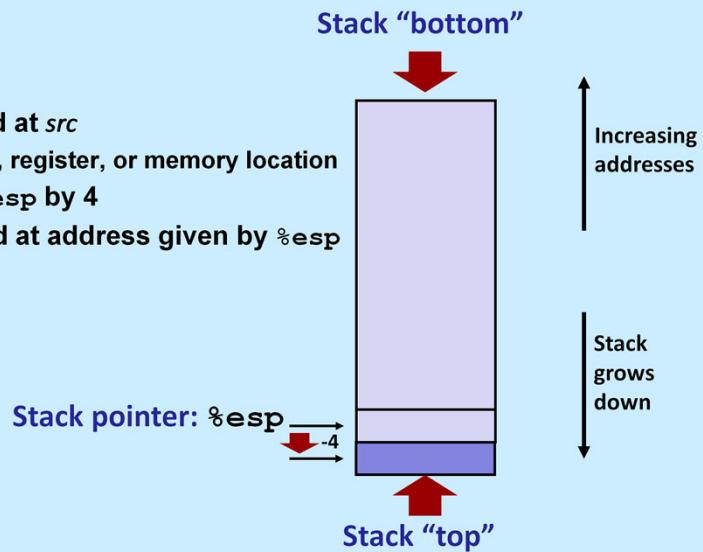
- Region of memory managed “last-in, first-out”
- Grows toward lower addresses
- Register `%esp` contains lowest stack address
 - address of “top” element



Supplied by CMU.

IA32 Stack: Push

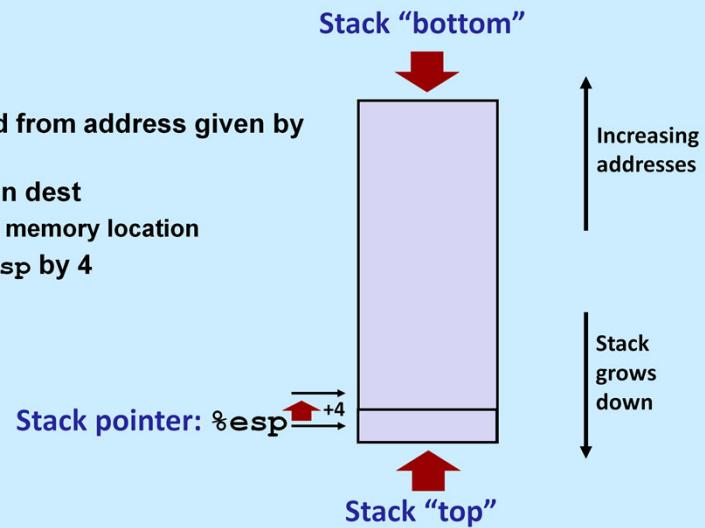
- **pushl src**
 - fetch operand at *src*
 - » immediate, register, or memory location
 - decrement %esp by 4
 - store operand at address given by %esp



Supplied by CMU.

IA32 Stack: Pop

- **popl dest**
 - fetch operand from address given by `%esp`
 - put operand in `dest`
 - » register or memory location
 - increment `%esp` by 4



Supplied by CMU.

Procedure Control Flow

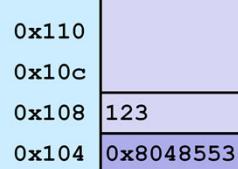
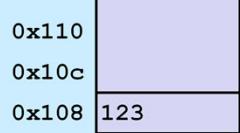
- Use stack to support procedure call and return
 - **Procedure call:** `call sub`
 - push return address on stack
 - jump to *sub*
 - **Return address:**
 - address of the next instruction after call
 - example from disassembly
- ```
804854e: e8 3d 06 00 00 call 8048b90 <sub>
8048553: 50 pushl %eax
```
- return address = 0x8048553
  - **Procedure return:** `ret`
    - pop address from stack
    - jump to address

Supplied by CMU.

## Procedure Call

```
804854e: e8 3d 06 00 00 call 8048b90 <sub>
8048553: 50 pushl %eax
```

**call 8048b90**



%esp 0x108

%esp 0x104

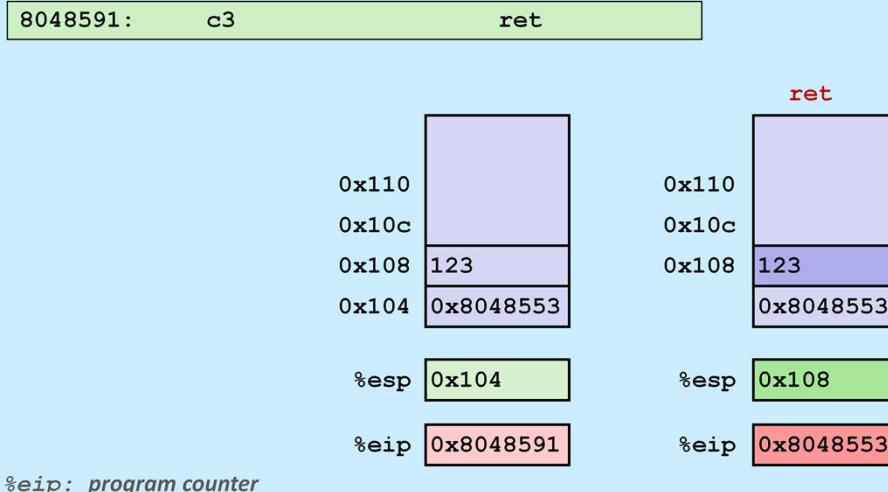
%eip 0x804854e

%eip 0x8048b90

*%eip: program counter*

Supplied by CMU.

## Procedure Return



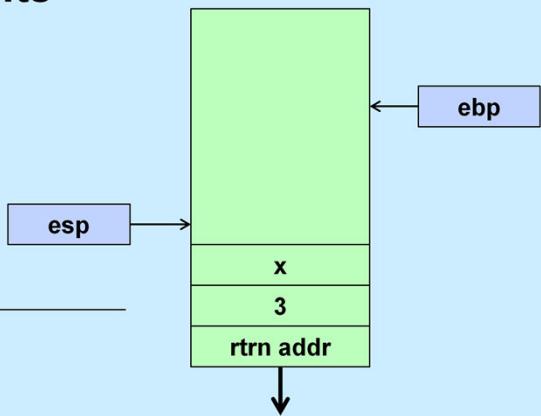
Supplied by CMU.

## Passing Arguments

```
int x;
int res;
int main() {
 ...
 res = subr(3, x);
 ...
}
```

---

```
main:
 ...
 pushl x
 pushl $3
 call subr
 movl %eax,res
 ...
```



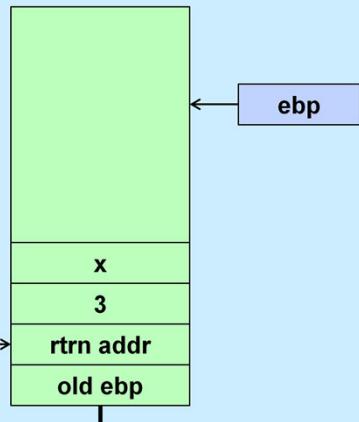
## Retrieving Arguments

```
int subr(int a, int b) {
 return a + b;
}
```

---

```
subr:
 pushl %ebp
 movl %esp, %ebp
 movl 12(%ebp), %eax
 addl 8(%ebp), %eax
 popl %ebp
 ret
```

---



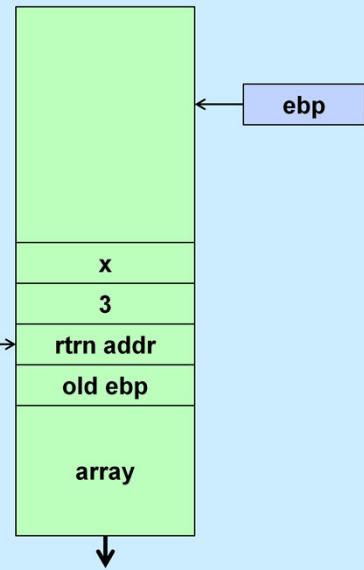
## Space for Local Variables

```
int subr(int a, int b) {
 int array[20];
 ...
}
```

---

```
subr:
 pushl %ebp
 movl %esp, %ebp
 subl $80, $esp
 ...
 addl $80, $esp
 popl %ebp
 ret
```

---



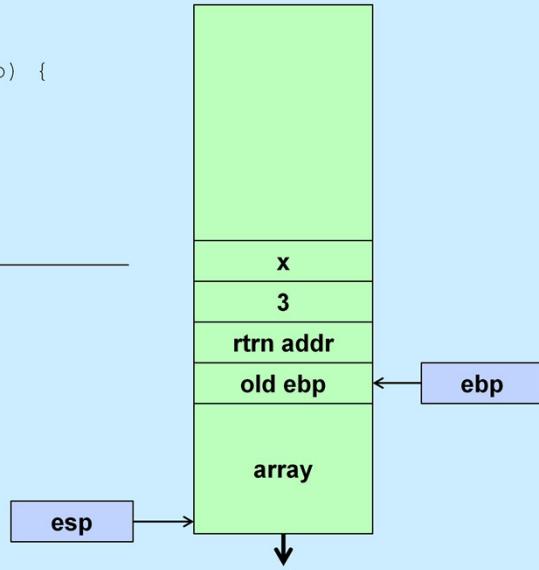
## Quick Exit ...

```
int subr(int a, int b) {
 int array[20];
 ...
}
```

---

```
subr:
 pushl %ebp
 movl %esp, %ebp
 subl $80, $esp
 ...
 leave
 ret
```

---



The *leave* instruction causes the contents of *ebp* to be copied into *esp*, thereby removing everything from the stack that had been pushed into the frame. It then pops the current stack top (the old *ebp*) into the *ebp* register. The effect of *leave* is thus to return to the caller's stack frame.

There is an *enter* instruction that has the same effect as that of the first three instructions of *subr* combined (it has an operand that indicates how much space for local variables to allocate). However, it's not used by *gcc*, apparently because it's slower than doing it as shown in the slide.

## Register-Saving Conventions

- When procedure `yoo` calls `who`:
  - `yoo` is the *caller*
  - `who` is the *callee*
- Can registers be used for temporary storage?

```
yoo:
 . . .
 movl $33, %edx
 call who
 addl %edx, %eax
 . . .
 ret
```

```
who:
 . . .
 movl 8(%ebp), %edx
 addl $32, %edx
 . . .
 ret
```

- contents of register `%edx` overwritten by `who`
- this could be trouble: something should be done!
  - » need some coordination

Supplied by CMU.

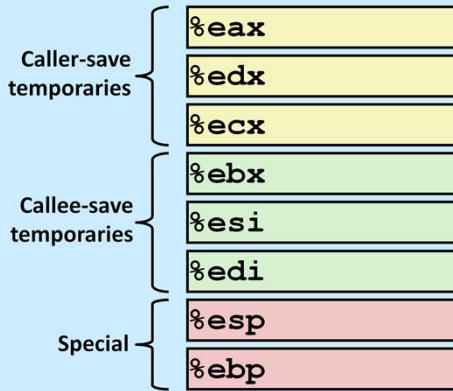
## Register-Saving Conventions

- When procedure `yoo` calls who:
  - `yoo` is the *caller*
  - `who` is the *callee*
- Can registers be used for temporary storage?
- Conventions
  - “*caller save*”
    - » caller saves temporary values on stack before the call
    - » restores them after call
  - “*callee save*”
    - » callee saves temporary values on stack before using
    - » restores them before returning

Supplied by CMU.

# IA32/Linux+Windows Register Usage

- %eax, %edx, %ecx
  - caller saves prior to call if values are used later
- %eax
  - also used to return integer value
- %ebx, %esi, %edi
  - callee saves if wants to use them
- %esp, %ebp
  - special form of callee-save
  - restored to original values upon exit from procedure



Supplied by CMU.

## Register-Saving Example

```
yoo:
...
 movl $33, %edx
 pushl %edx
 call who
 popl %edx
 addl %edx, %eax
...
 ret
```

```
who:
...
 pushl %ebx
...
 movl 4(%ebp), %ebx
 addl %53, %ebx
 movl 8(%ebp), %edx
 addl $32, %edx
...
 popl %ebx
...
 ret
```

Supplied by CMU.

## Recursive Function

```
/* Recursive popcount */
int pcount_r(unsigned x) {
 if (x == 0)
 return 0;
 else return
 (x & 1) + pcount_r(x >> 1);
}
```

- **Registers**

- **%eax, %edx used without first saving**
- **%ebx used, but saved at beginning & restored at end**

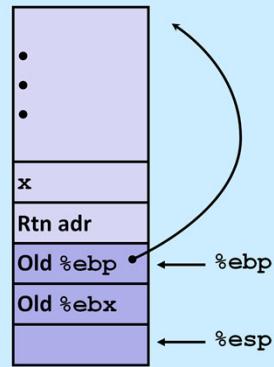
```
pcount_r:
 pushl %ebp
 movl %esp, %ebp
 pushl %ebx
 subl $4, %esp
 movl 8(%ebp), %ebx
 movl $0, %eax
 testl %ebx, %ebx
 je .L3
 movl %ebx, %eax
 shr1 %eax
 movl %eax, (%esp)
 call pcount_r
 movl %ebx, %edx
 andl $1, %edx
 leal (%edx,%eax), %eax
.L3:
 addl $4, %esp
 popl %ebx
 popl %ebp
 ret
```

Supplied by CMU.

## Recursive Call #1

```
/* Recursive popcount */
int pcount_r(unsigned x) {
 if (x == 0)
 return 0;
 else return
 (x & 1) + pcount_r(x >> 1);
}
```

```
pcount_r:
 pushl %ebp
 movl%esp, %ebp
 pushl %ebx
 subl$4, %esp
 movl8(%ebp), %ebx
 • • •
```



Supplied by CMU.

## Recursive Call #2

```
/* Recursive popcount */
int pcount_r(unsigned x) {
 if (x == 0)
 return 0;
 else return
 (x & 1) + pcount_r(x >> 1);
}
```

```
...
 movl $0, %eax
 testl %ebx, %ebx
 je .L3
 ...
.L3:
 ...
 ret
```

- Actions
  - if  $x == 0$ , return
    - » with  $%eax$  set to 0

$\%ebx$  x

Supplied by CMU.

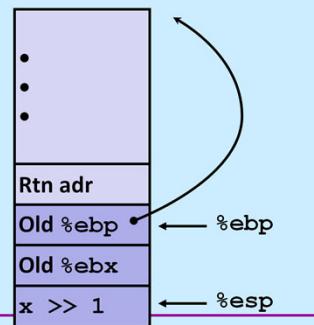
## Recursive Call #3

```
/* Recursive popcount */
int pcount_r(unsigned x) {
 if (x == 0)
 return 0;
 else return
 (x & 1) + pcount_r(x >> 1);
}
```

```
...
movl %ebx, %eax
shrl %eax
movl %eax, (%esp)
call pcount_r
...
```

- **Actions**
  - store  $x \gg 1$  on stack
  - make recursive call
- **Effect**
  - %eax set to function result
  - %ebx still has value of x

%ebx x



Supplied by CMU.

## Recursive Call #4

```
/* Recursive popcount */
int pcount_r(unsigned x) {
 if (x == 0)
 return 0;
 else return
 (x & 1) + pcount_r(x >> 1);
}

 • • •
 movl %ebx, %edx
 andl $1, %edx
 leal (%edx,%eax), %eax
 • • •
```

- **Assume**
  - %eax holds value from recursive call
  - %ebx holds x
- **Actions**
  - compute  $(x \& 1)$  + computed value
- **Effect**
  - %eax set to function result

%ebx x

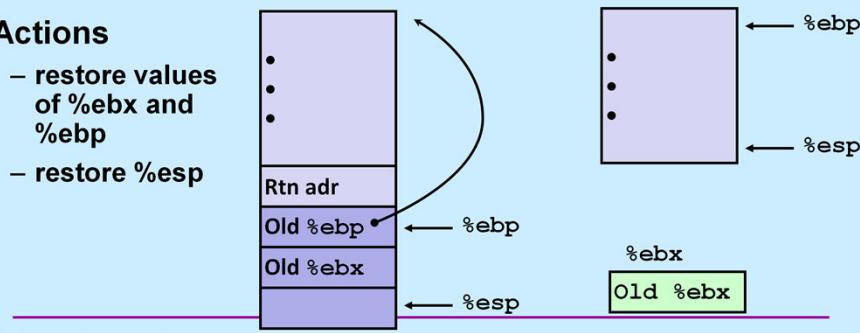
Supplied by CMU.

## Recursive Call #5

```
/* Recursive popcount */
int pcount_r(unsigned x) {
 if (x == 0)
 return 0;
 else return
 (x & 1) + pcount_r(x >> 1);
}
```

```
• • •
L3:
 addl$4, %esp
 popl%ebx
 popl%ebp
 ret
```

- Actions
  - restore values of %ebx and %ebp
  - restore %esp



Supplied by CMU.

## Observations About Recursion

- Handled without special consideration
  - stack frames mean that each function call has private storage
    - » saved registers & local variables
    - » saved return pointer
  - register-saving conventions prevent one function call from corrupting another's data
  - stack discipline follows call / return pattern
    - » if P calls Q, then Q returns before P
    - » last-in, first-out
- Also works for mutual recursion
  - P calls Q; Q calls P

Supplied by CMU.

# Pointer Code

## Generating Pointer

```
/* Compute x + 3 */
int add3(int x) {
 int localx = x;
 incr(&localx, 3);
 return localx;
}
```

## Referencing Pointer

```
/* Increment value by k */
void incr(int *ip, int k) {
 *ip += k;
}
```

- **add3 creates pointer and passes it to incr**

Supplied by CMU.

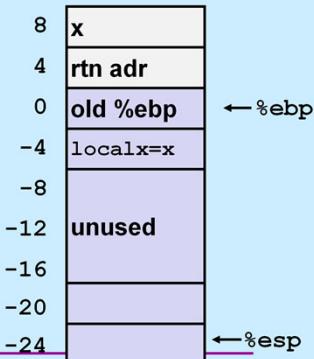
## Creating and Initializing Local Variables

```
int add3(int x) {
 int localx = x;
 incr(&localx, 3);
 return localx;
}
```

- Variable **localx** must be stored on stack
  - because: need to create pointer to it
- Compute pointer as **-4(%ebp)**

### First part of add3

```
add3:
 pushl %ebp
 movl %esp, %ebp
 subl $24, %esp # Alloc. 24 bytes
 movl 8(%ebp), %eax
 movl %eax, -4(%ebp) # Set localx to x
```



Supplied by CMU.

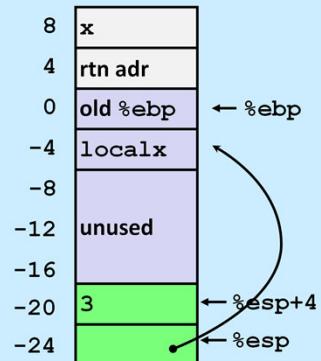
## Creating Pointer as Argument

```
int add3(int x) {
 int localx = x;
 incr(&localx, 3);
 return localx;
}
```

- Use **leal** instruction to compute address of **localx**

### Middle part of add3

```
movl $3, 4(%esp) # 2nd arg = 3
leal -4(%ebp), %eax# &localx
movl %eax, (%esp) # 1st arg = &localx
call incr
```



Supplied by CMU.

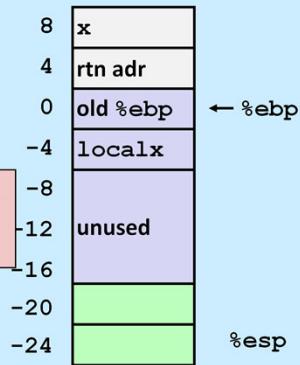
## Retrieving local variable

```
int add3(int x) {
 int localx = x;
 incr(&localx, 3);
 return localx;
}
```

- Retrieve localx from stack as return value

Final part of add3

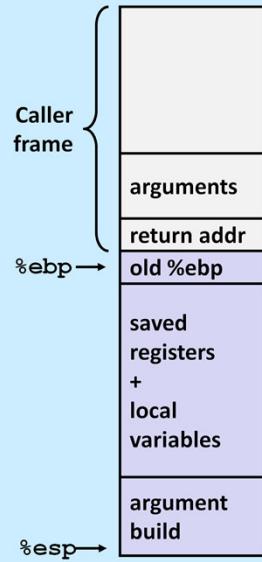
```
movl -4(%ebp), %eax # Return val= localx
leave
ret
```



Supplied by CMU.

# IA 32 Procedure Summary

- **Important Points**
  - stack is the right data structure for procedure call / return
    - » if P calls Q, then Q returns before P
- **Recursion (& mutual recursion) handled by normal calling conventions**
  - can safely store values in local stack frame and in callee-saved registers
  - put function arguments at top of stack
  - result return in %eax
- **Pointers are addresses of values**
  - on stack or global



Supplied by CMU.

## Why Bother with a Frame Pointer?

- It points to the beginning of the stack frame
  - making it easy for people to figure out where things are in the frame
  - but people don't execute the code ...
- The stack pointer always points somewhere within the stack frame
  - it moves about, but the compiler knows where it is pointing
    - » a local variable might be at 8(%rsp) for one instruction, but at 16(%rsp) for a subsequent one
    - » tough for people, but easy for the compiler
- Thus the frame pointer is superfluous
  - it can be used as a general-purpose register

## x86-64 General-Purpose Registers: Usage Conventions

|      |               |
|------|---------------|
| %rax | Return value  |
| %rbx | Callee saved  |
| %rcx | Argument #4   |
| %rdx | Argument #3   |
| %rsi | Argument #2   |
| %rdi | Argument #1   |
| %rsp | Stack pointer |
| %rbp | Callee saved  |
| %r8  | Argument #5   |
| %r9  | Argument #6   |
| %r10 | Caller saved  |
| %r11 | Caller Saved  |
| %r12 | Callee saved  |
| %r13 | Callee saved  |
| %r14 | Callee saved  |
| %r15 | Callee saved  |

Supplied by CMU.

## x86-64 Registers

- Arguments passed to functions via registers
  - if more than 6 integral parameters, then pass rest on stack
  - these registers can be used as caller-saved as well
- All references to stack frame via stack pointer
  - eliminates need to update %ebp/%rbp
- Other registers
  - 6 callee-saved
  - 2 caller-saved
  - 1 return value (also usable as caller-saved)
  - 1 special (stack pointer)

Supplied by CMU.

Note that the *leave* instruction is no longer relevant, since %rbp does not contain the address of the stack frame.

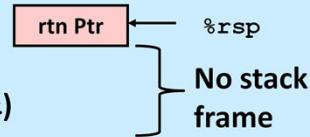
Also note that the conventions shown in the slide are those adopted by gcc on Linux; they aren't necessarily used by other compilers or on other operating systems. Even gcc doesn't use these conventions if optimization is completely turned off (in which case arguments are passed on the stack, just as for IA32).

## x86-64 Long Swap

```
void swap_l(long *xp, long *yp)
{
 long t0 = *xp;
 long t1 = *yp;
 *xp = t1;
 *yp = t0;
}
```

```
swap:
 movq (%rdi), %rdx
 movq (%rsi), %rax
 movq %rax, (%rdi)
 movq %rdx, (%rsi)
 ret
```

- Operands passed in registers
  - first (`xp`) in `%rdi`, second (`yp`) in `%rsi`
  - 64-bit pointers
- No stack operations required (except `ret`)
- Avoiding stack
  - can hold all local information in registers

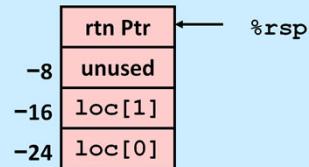


## x86-64 Locals in the Red Zone

```
/* Swap, using local array */
void swap_a(long *xp, long *yp)
{
 volatile long loc[2];
 loc[0] = *xp;
 loc[1] = *yp;
 *xp = loc[1];
 *yp = loc[0];
}
```

```
swap_a:
 movq (%rdi), %rax
 movq %rax, -24(%rsp)
 movq (%rsi), %rax
 movq %rax, -16(%rsp)
 movq -16(%rsp), %rax
 movq %rax, (%rdi)
 movq -24(%rsp), %rax
 movq %rax, (%rsi)
 ret
```

- **Avoiding stack-pointer change**
  - can hold all information within small window beyond stack pointer
    - » 128 bytes



Supplied by CMU.

The `volatile` keyword tells the compiler that it may not perform optimizations on the associated variable, such as storing it strictly in registers and not in memory. It's used primarily in cases where the variable might be modified via other routines that aren't apparent when the current code is being compiled. We'll see useful examples of its use later. Here it's used simply to ensure that `loc` is allocated on the stack, thus giving us a simple example of using local variables stored on the stack.

The issue here is whether a reference to memory beyond the current stack (as delineated by the stack pointer) is a legal reference. On IA32 it is not, but on x86-64 it is, as long as the reference is not more than 128 bytes beyond the end of the stack.

## x86-64 NonLeaf without Stack Frame

```
/* Swap a[i] & a[i+1] */
void swap_ele(long a[], int i)
{
 swap(&a[i], &a[i+1]);
}
```

- No values held while swap being invoked
- No callee-save registers needed
- rep instruction inserted as no-op
  - based on recommendation from AMD  
» can't handle transfer of control to ret

```
swap_ele:
 movslq %esi,%rsi # Sign extend i
 leaq 8(%rdi,%rsi,8), %rax # &a[i+1]
 leaq (%rdi,%rsi,8), %rdi # &a[i] (1st arg)
 movq %rax, %rsi # (2nd arg)
 call swap
 rep # No-op
 ret
```

Supplied by CMU.

## x86-64 Stack Frame Example

```
long sum = 0;
/* Swap a[i] & a[i+1] */
void swap_ele_su
 (long a[], int i)
{
 swap(&a[i], &a[i+1]);
 sum += (a[i]*a[i+1]);
}
```

- Keeps values of `&a[i]` and `&a[i+1]` in callee-save registers
  - `rbx` and `rbp`
- Must set up stack frame to save these registers
  - else clobbered in swap

```
swap_ele_su:
 movq %rbx, -16(%rsp)
 movq %rbp, -8(%rsp)
 subq $16, %rsp
 movslq %esi,%rax
 leaq 8(%rdi,%rax,8), %rbx
 leaq (%rdi,%rax,8), %rbp
 movq %rbx, %rsi
 movq %rbp, %rdi
 call swap
 movq (%rbx), %rax
 imulq (%rbp), %rax
 addq %rax, sum(%rip)
 movq (%rsp), %rbx
 movq 8(%rsp), %rbp
 addq $16, %rsp
 ret
```

Supplied by CMU.

Note that `sum` is a global variable. While its exact location in memory is not known by the compiler, it will be stored in memory at some location just beyond the end of the executable code (which is known as “text”). Thus the compiler can refer to `sum` via the instruction pointer. The actual displacement, i.e., the distance from the current target of the instruction pointer and the location of `sum`, is not known to the compiler, but will be known to the linker, which will fill this displacement in when the program is linked. This will all be explained in detail in two or three weeks.

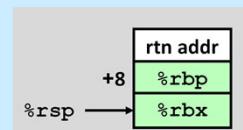
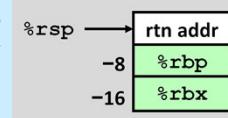
## Understanding x86-64 Stack Frame

```
swap_ele_su:
 movq %rbx, -16(%rsp) # Save %rbx
 movq %rbp, -8(%rsp) # Save %rbp
 subq $16, %rsp # Allocate stack frame
 movslq %esi,%rax # Extend i
 leaq 8(%rdi,%rax,8), %rbx # &a[i+1] (callee save)
 leaq (%rdi,%rax,8), %rbp # &a[i] (callee save)
 movq %rbx, %rsi # 2nd argument
 movq %rbp, %rdi # 1st argument
 call swap
 movq (%rbx), %rax # Get a[i+1]
 imulq (%rbp), %rax # Multiply by a[i]
 addq %rax, sum(%rip) # Add to sum
 movq (%rsp), %rbx # Restore %rbx
 movq 8(%rsp), %rbp # Restore %rbp
 addq $16, %rsp # Deallocate frame
 ret
```

Supplied by CMU.

## Understanding x86-64 Stack Frame

```
movq %rbx, -16(%rsp) # Save %rbx
movq %rbp, -8(%rsp) # Save %rbp
subq $16, %rsp # Allocate stack frame
• • •
movq (%rsp), %rbx # Restore %rbx
movq 8(%rsp), %rbp # Restore %rbp
addq $16, %rsp # Deallocate frame
```



Supplied by CMU.

## Interesting Features of Stack Frame

- **Allocate entire frame at once**
  - all stack accesses can be relative to `%rsp`
  - do by decrementing stack pointer
  - can delay allocation, since safe to temporarily use red zone
- **Simple deallocation**
  - increment stack pointer
  - no base/frame pointer needed

Supplied by CMU.

## x86-64 Procedure Summary

- Heavy use of registers
  - parameter passing
  - more temporaries since more registers
- Minimal use of stack
  - sometimes none
  - allocate/deallocate entire block
- Many tricky optimizations
  - what kind of stack frame to use
  - various allocation techniques

Supplied by CMU.

## Tail Recursion

```
int factorial(int x) { int factorial(int x) {
 if (x == 1) return f2(x, 1);
 return x; }
 else
 return
 x*factorial(x-1);
}

int f2(int a1, int a2) {
 if (a1 == 1)
 return a2;
 else
 return
 f2(a1-1, a1*a2);
}
```

The slide shows two implementations of the factorial function. Both use recursion. In the version on the left, the result of each recursive call is used within the invocation that issued the call. In the second, the result of each recursive call is simply returned. This is known as *tail recursion*.

## No Tail Recursion (1)

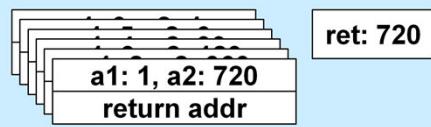
|             |
|-------------|
| x: 6        |
| return addr |
| x: 5        |
| return addr |
| x: 4        |
| return addr |
| x: 3        |
| return addr |
| x: 2        |
| return addr |
| x: 1        |
| return addr |

Here we look at the stack usage for the version without tail recursion. Note that we have as many stack frames as the value of the argument; the results of the calls are combined after the stack reaches its maximum size.

## No Tail Recursion (2)

|             |          |
|-------------|----------|
| x: 6        | ret: 720 |
| return addr |          |
| x: 5        | ret: 120 |
| return addr |          |
| x: 4        | ret: 24  |
| return addr |          |
| x: 3        | ret: 6   |
| return addr |          |
| x: 2        | ret: 2   |
| return addr |          |
| x: 1        | ret: 1   |
| return addr |          |

# Tail Recursion



With tail recursion, since the result of the recursive call is not used by the issuing stack frame, it's possible to reuse the issuing stack frame to handle the recursive invocation. Thus rather than push a new stack frame on the stack, the current one is written over. Thus the entire sequence of recursive calls can be handled within a single stack frame.

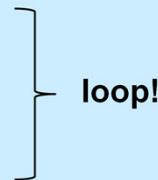
## Code: gcc -O1

```
f2:
 movl %esi, %eax
 cmpl $1, %edi
 je .L5
 subq $8, %rsp
 movl %edi, %esi
 imull %eax, %esi
 subl $1, %edi
 call f2 # recursive call!
 addq $8, %rsp
.L5:
 rep
 ret
```

This is the result of compiling the tail-recursive version of factorial using gcc with the `-O1` flag. This flag turns on a moderate level of code optimization, but not enough to cause the stack frame to be reused.

## Code: gcc -O2

```
f2:
 cmpl $1, %edi
 movl %esi, %eax
 je .L8
.L12:
 imull %edi, %eax
 subl $1, %edi
 cmpl $1, %edi
 jne .L12
.L8:
 rep
 ret
```



**loop!**

Here we've compiled the program using the `-O2` flag, which turns on additional optimization (at the cost of increased compile time), with the result that the recursive calls are optimized away — they are replaced with a loop.

Why not always compile with `-O2`? For “production code” that is fully debugged (assuming this is possible), this is a good idea. But this and other aggressive optimizations make it difficult to relate the runtime code with the source code. Thus, a runtime error might occur at some point in the program’s execution, but it is impossible to determine exactly which line of the source code was in play when the error occurred.