

inst.eecs.berkeley.edu/~cs61c
CS61C : Machine Structures

Lecture 18 – MapReduce


2014-10-13

Senior Lecturer SOE Dan Garcia
www.cs.berkeley.edu/~ddgarcia

Top 10 Tech Trends ⇒

(1) Computing everywhere
 (2) Internet of Things (3) 3D printing (4) Analytics everywhere (5) Context rich systems (6) smart machines (7) cloud computing (8) software applications (9) web-scale IT (10) Security. Agree?

www.computerworld.com/article/2692619/gartner-lays-out-its-top-10-tech-trends-for-2015.html



Review of Last Lecture

- Warehouse Scale Computing
 - Example of parallel processing in the post-PC era
 - Servers on a rack, rack part of cluster
 - Issues to handle include **load balancing**, **failures**, **power usage** (sensitive to cost & energy efficiency)
- PUE = Total building power / IT equipment power

2

Great Idea #4: Parallelism

Today's Lecture

Software

- Parallel Requests
Assigned to computer
e.g. Search "Garcia"
- Parallel Threads
Assigned to core
e.g. Lookup, Ads
- Parallel Instructions
> 1 instruction @ one time
e.g. 5 pipelined instructions
- Parallel Data
> 1 data item @ one time
e.g. add of 4 pairs of words
- Hardware descriptions
All gates functioning in parallel at same time

Hardware

Warehouse Scale Computer

Smart Phone

Leverage Parallelism & Achieve High Performance

Computer

Core ... Core

Memory

Input/Output

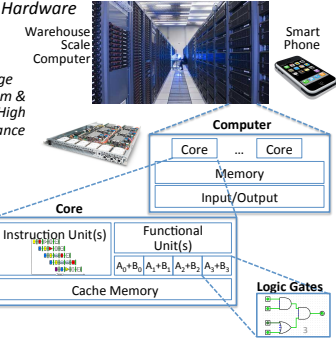
Core

Instruction Unit(s)

Functional Unit(s)

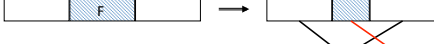
Cache Memory

Logic Gates



Amdahl's (Heartbreaking) Law

- Speedup due to enhancement E:

$$\text{Speedup } w/E = \frac{\text{Exec time w/o E}}{\text{Exec time w/E}}$$
 - Example:** Suppose that enhancement E accelerates a fraction F (F<1) of the task by a factor S (S>1) and the remainder of the task is unaffected
- 
- Exec time w/E = Exec Time w/o E × [(1-F) + F/S]
 Speedup w/E = 1 / [(1-F) + F/S]

4

Amdahl's Law


- Speedup =
$$\frac{1}{(1 - F) + \frac{F}{S}}$$

Non-spud-up part → (1 - F) F/S → Spud-up part
 - Example:** the execution time of 1/5 of the program can be accelerated by a factor of 10. What is the program speed-up overall?
- $$\frac{1}{0.8 + \frac{0.2}{10}} = \frac{1}{0.8 + 0.02} = 1.22$$

5

Consequence of Amdahl's Law

- The amount of speedup that can be achieved through parallelism is limited by the non-parallel portion of your program!



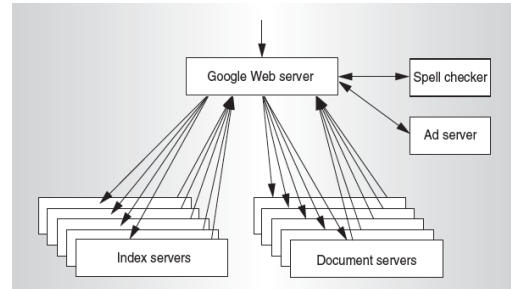
6

Request-Level Parallelism (RLP)

- Hundreds or thousands of requests per sec
 - Not your laptop or cell-phone, but popular Internet services like web search, social networking, ...
 - Such requests are largely independent
 - Often involve read-mostly databases
 - Rarely involve strict read-write data sharing or synchronization across requests
- Computation easily partitioned within a request and across different requests

7

Google Query-Serving Architecture



8

Data-Level Parallelism (DLP)

- Two kinds:
 - 1) Lots of **data in memory** that can be operated on in parallel (e.g. adding together 2 arrays)
 - 2) Lots of **data on many disks** that can be operated on in parallel (e.g. searching for documents)
- 1) SIMD does Data-Level Parallelism (DLP) in memory
- 2) Today's lecture, Lab 6, Proj. 3 do DLP across many servers and disks using **MapReduce**

9

Administrivia ... The Midterm

- Average around 10/20
 - Despite lots of partial credit
 - Regrades being processed
 - Have perspective – it's only 20 / 300 points.
 - Don't panic. Do lots of practice problems in a team. Do NOT study alone.
- Part 2 will be easier to compensate
- You can clobber Part 1 with Part 2

10

What is MapReduce?

- Simple data-parallel programming model designed for scalability and fault-tolerance
- Pioneered by Google
 - Processes > 25 petabytes of data per day
- Popularized by open-source Hadoop project
 - Used at Yahoo!, Facebook, Amazon, ...



11

What is MapReduce used for?

- At Google:
 - Index construction for Google Search
 - Article clustering for Google News
 - Statistical machine translation
 - For computing multi-layer street maps
- At Yahoo!:
 - "Web map" powering Yahoo! Search
 - Spam detection for Yahoo! Mail
- At Facebook:
 - Data mining
 - Ad optimization
 - Spam detection

12

Example: Facebook Lexicon



13


MapReduce Design Goals

1. Scalability to large data volumes:
 - 1000's of machines, 10,000's of disks
2. Cost-efficiency:
 - Commodity machines (cheap, but unreliable)
 - Commodity network
 - Automatic fault-tolerance via re-execution (fewer administrators)
 - Easy, fun to use (fewer programmers)

Jeffrey Dean and Sanjay Ghemawat, "MapReduce: Simplified Data Processing on Large Clusters," 6th USENIX Symposium on Operating Systems Design and Implementation, 2004. (optional reading, linked on course homepage – a digestible CS paper at the 61C level)

14

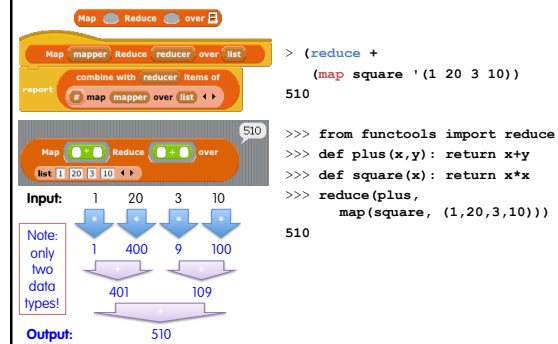
Typical Hadoop Cluster



- 40 nodes/rack, 1000-4000 nodes in cluster
- 1 Gbps bandwidth within rack, 8 Gbps out of rack
- Node specs (Yahoo terasort):
8 x 2GHz cores, 8 GB RAM, 4 disks (= 4 TB?)

15

MapReduce in CS10 & CS61A{,S}



MapReduce Programming Model

Input & Output: each a set of key/value pairs

Programmer specifies two functions:

map (**in_key**, **in_value**) → **list**(**interm_key**, **interm_value**)

- Processes input key/value pair
- Slices data into "shards" or "splits"; distributed to workers
- Produces set of intermediate pairs

reduce (**interm_key**, **list**(**interm_value**)) → **list**(**out_value**)

- Combines all intermediate values for a particular key
- Produces a set of merged output values (usu just one)

code.google.com/edu/parallel/mapreduce-tutorial.html

MapReduce WordCount Example

- "Mapper" nodes are responsible for the **map** function

```

// "I do I learn" ⇒ ("I",1), ("do",1), ("I",1), ("learn",1)
map(String input_key,
  String input_value):
  // input_key : document name (or line of text)
  // input_value: document contents
  for each word w in input_value:
    EmitIntermediate(w, "1");
  
```

- "Reducer" nodes are responsible for the **reduce** function

```

// ("I",[1,1]) ⇒ ("I",2)
reduce(String output_key,
  Iterator intermediate_values):
  // output_key : a word
  // output_values: a list of counts
  int result = 0;
  for each v in intermediate_values:
    result += ParseInt(v);
  Emit(AsString(result));
  
```

- Data on a distributed file system (DFS)

